

What follows is the article “A Framework for Building Verified Partial Evaluators” authored by Jason Gross (myself), Adam Chlipala, and Andres Erbsen, submitted to PLDI’20, the ACM SIGPLAN Conference on Programming Language Design and Implementation, which will take place June 15–20, 2020 in London, United Kingdom.

Since my contribution is not clearly identified in the conference paper, I describe the distribution of authorship here. Nearly all of the implementation and proof work of the partial evaluator, the artifact described in this paper, was done by me; others have co-authored some of the common utility files (such as the library on facts about binary integers), and much of arithmetic templates of Fiat Cryptography, used as the main case-study for this paper, were written by Andres Erbsen and Jade Philipoom. The design of the rewriter was hashed out with help from Andres Erbsen, Jade Philipoom, and Adam Chlipala over the course of many discussions. All performance tests and aggregation of results were done by me. I wrote most of the technical content of the paper; Adam contributed technical explanations of some of the prior work. Adam also organized most of the paper and wrote most of the introduction and related works sections, as well as most of the prose gluing the various paragraphs about technical content together, with some help from me and Andres.

# A Framework for Building Verified Partial Evaluators

Anonymous Author(s)

## Abstract

*Partial evaluation* is a classic technique for generating lean, customized code from libraries that start with more bells and whistles. It is also an attractive approach to creation of *formally verified* systems, where theorems can be proved about libraries, yielding correctness of all specializations “for free.” However, it can be challenging to make library specialization both performant and trustworthy. We present a new approach, prototyped in the Coq proof assistant, which supports specialization at the speed of native-code execution, without adding to the trusted code base. Our extensible engine, which combines the traditional concepts of tailored term reduction and automatic rewriting from hint databases, is also of interest to replace these ingredients in proof assistants’ proof checkers and tactic engines, at the same time as it supports extraction to standalone compilers from library parameters to specialized code.

## 1 Introduction

Mechanized proof is gaining in importance for development of critical software infrastructure. Oft-cited examples include the CompCert verified C compiler [17] and the seL4 verified operating-system microkernel [16]. Here we have very flexible systems that are ready to adapt to varieties of workloads, be they C source programs for CompCert or application binaries for seL4. For a verified operating system, such adaptation takes place at *runtime*, when we launch the application. However, some important bits of software infrastructure commonly do adaptation at *compile time*, such that the fully general infrastructure software is not even installed in a deployed system.

Of course, compilers are a natural example of that pattern, as we would not expect CompCert itself to be installed on an embedded system whose application code was compiled with it. The problem is that writing a compiler is rather labor-intensive, with its crafting of syntax-tree types for source, target, and intermediate languages, its fine-tuning of code for transformation passes that manipulate syntax trees explicitly, and so on. An appealing alternative is *partial evaluation* [15], which relies on reusable compiler facilities to specialize library code to parameters, with no need to write that library code in terms of syntax-tree manipulations. Cutting-edge tools in this tradition even make it possible to

use high-level functional languages to generate performance-competitive low-level code, as in Scala’s Lightweight Modular Staging [22].

It is natural to try to port this approach to construction of systems with mechanized proofs. On one hand, the typed functional languages in popular proof assistants’ logics make excellent hosts for flexible libraries, which can often be specialized through means as simple as partial application of curried functions. Term-reduction systems built into the proof assistants can then generate the lean residual programs. On the other hand, it is surprisingly difficult to realize the last sentence with good performance. The challenge is that we are not just implementing algorithms; we also want a proof to be checked by a small proof checker, and there is tension in designing such a checker, as fancier reduction strategies grow the trusted code base. It would seem like an abandonment of the spirit of proof assistants to bake in a reduction strategy per library, yet effective partial evaluation tends to be rather fine-tuned in this way. Performance tuning matters when generated code is thousands of lines long.

In this paper, we present an approach to verified partial evaluation in proof assistants, which requires no changes to proof checkers. To make the relevance concrete, we use the example of Fiat Cryptography [11], a Coq library that generates code for big-integer modular arithmetic at the heart of elliptic-curve cryptography algorithms. This domain-specific compiler has been adopted, for instance, in the Chrome Web browser, such that about half of all HTTPS connections from browsers are now initiated using code generated (with proof) by Fiat Cryptography. However, Fiat Cryptography was only used successfully to build C code for the two most widely used curves (P-256 and Curve25519). Their method of partial evaluation timed out trying to compile code for the third most widely used curve (P-384). Additionally, to achieve acceptable reduction performance, the library code had to be written manually in continuation-passing style. We will demonstrate a new Coq library that corrects both weaknesses, while maintaining the generality afforded by allowing rewrite rules to be mixed with partial evaluation.

### 1.1 A Motivating Example

We are interested in partial-evaluation examples that mix higher-order functions, inductive datatypes, and arithmetic simplification. For instance, consider the following Coq code.

```
Definition prefixSums (ls:list nat) : list nat :=  
  let ls' := combine ls (seq 0 (length ls)) in  
  let ls'' := map (λ p, fst p * snd p) ls' in  
  let '(_, ls''') := fold_left (λ acc_ls''' n,  
    let 'acc, ls'''') := acc_ls''' in
```

```

111  let acc' := acc + n in
112    (acc', acc' :: ls'') ls'' (0, [])
113    ls''.
114
115 This function first computes list ls' that pairs each element of input list ls with its position, so, for instance, list [a; b; c] becomes [(a, 0); (b, 1); (c, 2)]. Then we map over the list of pairs, multiplying the components at each position. Finally, we traverse that list, building up a list of all prefix sums.
116
117
118
119
120

```

We would like to specialize this function to particular list lengths. That is, we know in advance how many list elements we will pass in, but we do not know the values of those elements. For a given length, we can construct a schematic list with one free variable per element. For example, to specialize to length four, we can apply the function to list [a; b; c; d], and we expect this output:

```

121  let acc := b + c * 2 in
122  let acc' := acc + d * 3 in
123    [acc'; acc; b; 0]
124
125
126
127

```

Notice how subterm sharing via `lets` is important. As list length grows, we avoid quadratic blowup in term size through sharing. Also notice how we simplified the first two multiplications with  $a \cdot 0 = 0$  and  $b \cdot 1 = b$  (each of which requires explicit proof in Coq), using other arithmetic identities to avoid introducing new variables for the first two prefix sums of ls'', as they are themselves constants or variables, after simplification.

To set up our partial evaluator, we prove the algebraic laws that it should use for simplification, starting with basic arithmetic identities.

```

128 Lemma zero_plus : forall n, 0 + n = n.
129 Lemma plus_zero : forall n, n + 0 = n.
130 Lemma times_zero : forall n, n * 0 = 0.
131 Lemma times_one : forall n, n * 1 = n.
132
133
134
135
136
137
138
139
140
141
142
143
144
145
146
147
148
149
150
151
152
153
154
155
156
157
158
159
160
161
162
163
164
165

```

Next, we prove a law for each list-related function, connecting it to the primitive-recursion combinator for some inductive type (natural numbers or lists, as appropriate). We use a special apostrophe marker to indicate a quantified variable that may only match with *compile-time constants*. We also use a further marker `ident.eagerly` to ask the reducer to simplify a case of primitive recursion by complete traversal of the designated argument's constructor tree.

```

156 Lemma eval_map A B (f : A -> B) l
157   : map f l = ident.eagerly list_rect _ _ []
158   (lambda x _ l', f x :: l') l.
159 Lemma eval_fold_left A B (f : A -> B -> A) l a
160   : fold_left f l a = ident.eagerly list_rect
161   _ _ (lambda a, a)
162   (lambda x _ r a, r (f a x)) l a.
163 Lemma eval_combine A B (la : list A) (lb : list B)
164   : combine la lb = list_rect _ (lambda _, [])
165   (lambda x _ r lb, list_case (lambda _, _) []
166   (lambda y ys, (x, y) :: r ys) lb) la lb.
167
168
169
170
171
172
173
174
175
176
177
178
179
180
181
182
183
184
185
186
187
188
189
190
191
192
193
194
195
196
197
198
199
200
201
202
203
204
205
206
207
208
209
210
211
212
213
214
215
216
217
218
219
220

```

```
Lemma eval_length A (ls : list A)
```

```
: length ls = list_rect _ 0 (lambda _ _ n, S n) ls.
```

With all the lemmas available, we can package them up into a rewriter, which triggers generation of a specialized rewrite procedure and its soundness proof. Our Coq plugin introduces a new command `Make` for building rewriters

```
Make rewriter := Rewriter For (zero_plus, plus_zero,
times_zero, times_one, eval_map, eval_fold_left,
do_again eval_length, do_again eval_combine,
eval_rect nat, eval_rect list, eval_rect prod)
  (with delta) (with extra idents (seq)).
```

Most inputs to `Rewriter For` list quantified equalities to use for left-to-right rewriting. However, we also use options `do_again`, to request that some rules trigger an extra bottom-up pass after being used for rewriting; `eval_rect`, to queue up eager evaluation of a call to a primitive-recursion combinator on a known recursive argument; `with delta`, to request evaluation of all monomorphic operations on concrete inputs; and `with extra idents`, to inform the engine of further permitted identifiers that do not appear directly in any of the rewrite rules.

Our plugin also provides new tactics like `Rewrite_rhs_for`, which applies a rewriter to the righthand side of an equality goal. That last tactic is just what we need to synthesize a specialized `prefixSums` for list length four, along with a proof of its equivalence to the original function.

```
Definition prefixSums4 :
{f : nat -> nat -> nat -> nat -> list nat}
| forall a b c d, f a b c d = prefixSums [a;b;c;d]] :=
ltac:(eexists; Rewrite_rhs_for rewriter; reflexivity).
```

## 1.2 Concerns of Trusted-Code-Base Size

Crafting a reduction strategy is challenging enough in a standalone tool. A large part of the difficulty in a proof assistant is reducing in a way that leaves a proof trail that can be checked efficiently by a small kernel. Most proof assistants present user-friendly surface tactic languages that generate proof traces in terms of more elementary tactic steps. The trusted proof checker only needs to know about the elementary steps, and there is pressure to be sure that these steps are indeed elementary, not requiring excessive amounts of kernel code. However, hardcoding a new reduction strategy in the kernel can bring dramatic performance improvements. Generating thousands of lines of code with partial evaluation would be intractable if we were outputting sequences of primitive rewrite steps justifying every little term manipulation, so we must take advantage of the time-honored feature of type-theoretic proof assistants that reductions included in the definitional equality need not be requested explicitly.

Which kernel-level reductions *does* Coq support today? Currently, the trusted code base knows about four different kinds of reduction: left-to-right conversion, right-to-left conversion, a virtual machine (VM) written in C based on the OCaml compiler, and a compiler to native code. Furthermore,

the first two are parameterized on an arbitrary user-specified ordering of which constants to unfold when, in addition to internal heuristics about what to do when the user has not specified an unfolding order for given constants. Recently, native support for 63-bit integers has been added to the VM and native machines. A recent pull request proposes adding support for native IEEE 754-2008 binary64 floats [21], and support for native arrays is in the works [10].

To summarize, there has been quite a lot of “complexity creep” in the Coq trusted base, to support efficient reduction, and yet realistic partial evaluation has *still* been rather challenging. Even the additional three reduction mechanisms outside Coq’s kernel (cbn, simpl, cbv) are not at first glance sufficient for verified partial evaluation.

### 1.3 Our Solution

Aehlig et al. [1] presented a very relevant solution to a related problem, using *normalization by evaluation (NbE)* [4] to bootstrap reduction of open terms on top of full reduction, as built into a proof assistant. However, it was simultaneously true that they expanded the proof-assistant trusted code base in ways specific to their technique, and that they did not report any experiments actually using the tool for partial evaluation (just traditional full reduction), potentially hiding performance-scaling challenges or other practical issues. We have adapted their approach in a new Coq library embodying **the first partial-evaluation approach to satisfy the following criteria.**

- It integrates with a general-purpose, foundational proof assistant, **without growing the trusted base**.
- For a wide variety of initial functional programs, it provides **fast** partial evaluation with reasonable memory use.
- It allows reduction that **mixes rules of the definitional equality** with *equalities proven explicitly as theorems*.
- It **preserves sharing** of common subterms.
- It also allows **extraction of standalone partial evaluators**.

Our contributions include answers to a number of challenges that arise in scaling NbE-based partial evaluation in a proof assistant. First, we rework the approach of Aehlig et al. [1] to function *without extending a proof assistant’s trusted code base*, which, among other challenges, requires us to prove termination of reduction and encode pattern matching explicitly (leading us to adopt the performance-tuned approach of Maranget [20]).

Second, using partial evaluation to generate residual terms thousands of lines long raises *new scaling challenges*:

- Output terms may contain *so many nested variable binders* that we expect it to be performance-prohibitive to perform bookkeeping operations on first-order-encoded terms (e.g., with de Bruijn indices, as is done in  $\mathcal{R}_{\text{tac}}$  by Malecha and Bengtson [18]). For instance, while

the reported performance experiments of Aehlig et al. [1] generate only closed terms with no binders, Fiat Cryptography may generate a single routine (e.g., multiplication for curve P-384) with nearly a thousand nested binders.

- Naive representation of terms without proper *sharing of common subterms* can lead to fatal term-size blow-up. Fiat Cryptography’s arithmetic routines rely on significant sharing of this kind.
- Unconditional rewrite rules are in general insufficient, and we need *rules with side conditions*. For instance, in Fiat Cryptography, some rules for simplifying modular arithmetic depend on proofs that operations in subterms do not overflow.
- However, it is also not reasonable to expect a general engine to discharge all side conditions on the spot. We need integration with *abstract interpretation* that can analyze whole programs to support reduction.

Briefly, our respective solutions to these problems are the *parametric higher-order abstract syntax (PHOAS)* [8] term encoding, a *let-lifting* transformation threaded throughout reduction, extension of rewrite rules with executable Boolean side conditions, and a design pattern that uses decorator function calls to include analysis results in a program.

Finally, we carry out the *first large-scale performance-scaling evaluation* of partial evaluation in a proof assistant, covering all elliptic curves from the published Fiat Cryptography experiments, along with microbenchmarks.

This paper proceeds through explanations of the trust stories behind our approach and earlier ones (section 2), the core structure of our engine (section 3), the additional scaling challenges we faced (section 4), performance experiments (section 5), and related work (section 6) and conclusions. Our implementation is included as an anonymous supplement.

## 2 Trust, Reduction, and Rewriting

Since much of the narrative behind our design process depends on tradeoffs between performance and trustworthiness, we start by reviewing the general situation in proof assistants.

Across a variety of proof assistants, simplification of functional programs is a workhorse operation. Proof assistants like Coq that are based on type theory typically build in *definitional equality* relations, identifying terms up to reductions like  $\beta$ -reduction and unfolding of named identifiers. What looks like a single “obvious” step in an on-paper equational proof may require many of these reductions, so it is handy to have built-in support for checking a claimed reduction. Figure 1a diagrams how such steps work in a system like Coq, where the system implementation is divided between a trusted *kernel*, for checking *proof terms* in a minimal language, and additional untrusted support, like a *tactic engine*

evaluating a language of higher-level proof steps, in the process generating proof terms out of simpler building blocks. It is standard to include a primitive proof step that validates any reduction compatible with the definitional equality, as the latter is decidable. The figure shows a tactic that simplifies a goal using that facility.

In proof goals containing free variables, executing subterms can get stuck before reaching normal forms. However, we can often achieve further simplification by using equational rules that we prove explicitly, rather than just relying on the rules built into the definitional equality and its decidable equivalence checker. Coq's autorewrite tactic, as diagrammed in Figure 1b, is a good example: it takes in a database of quantified equalities and applies them repeatedly to rewrite in a goal. It is important that Coq's kernel does not trust the autorewrite tactic. Instead, the tactic must output a proof term that, in some sense, is the moral equivalent of a line-by-line equational proof. It can be challenging to keep these proof terms small enough, as naive rewrite-by-rewrite versions repeatedly copy large parts of proof goals, justifying a rewrite like  $C[e_1] = C[e_2]$  for some context  $C$  given a proof of  $e_1 = e_2$ , with the full value of  $C$  replicated in the proof term for that single rewrite. Overcoming these challenges while retaining decidability of proof checking is tricky, since we may use autorewrite with rule sets that do not always lead to terminating reduction. Coq includes more experimental alternatives like rewrite\_strat, which use bottom-up construction of multi-rewrite proofs, with sharing of common contexts. Still, as section 5 will show, these methods that generate substantial proof terms are at significant performance disadvantages.

Now we summarize how Aehlig et al. [1] provide flexible and fast interleaving of standard  $\lambda$ -calculus reduction and use of proved equalities (the next section will go into more detail). Figure 1c demonstrates a workflow based on *a deep embedding of a core ML-like language*. That is, within the logic of the proof assistant (Isabelle/HOL, in their case), a type of syntax trees for ML programs is defined, with an associated operational semantics. The basic strategy is, for a particular set of rewrite rules and a particular term to simplify, to generate a (deeply embedded) ML program that, if it terminates, produces a syntax tree for the simplified term. Their tactic uses reification to create ML versions of rule sets and terms. They also wrote a reduction function in ML and proved it sound once and for all, against the ML operational semantics. Combining that proof with proofs generated by reification, we conclude that an application of the reduction function to the reified rules and term is indeed an ML term that generates correct answers. The tactic then “throws the ML term over the wall,” using a general code-generation framework for Isabelle/HOL [14]. Trusted code compiles the ML code into the concrete syntax of a mainstream ML language, Standard ML in their case, and compiles it with an off-the-shelf compiler. The output of that compiled program

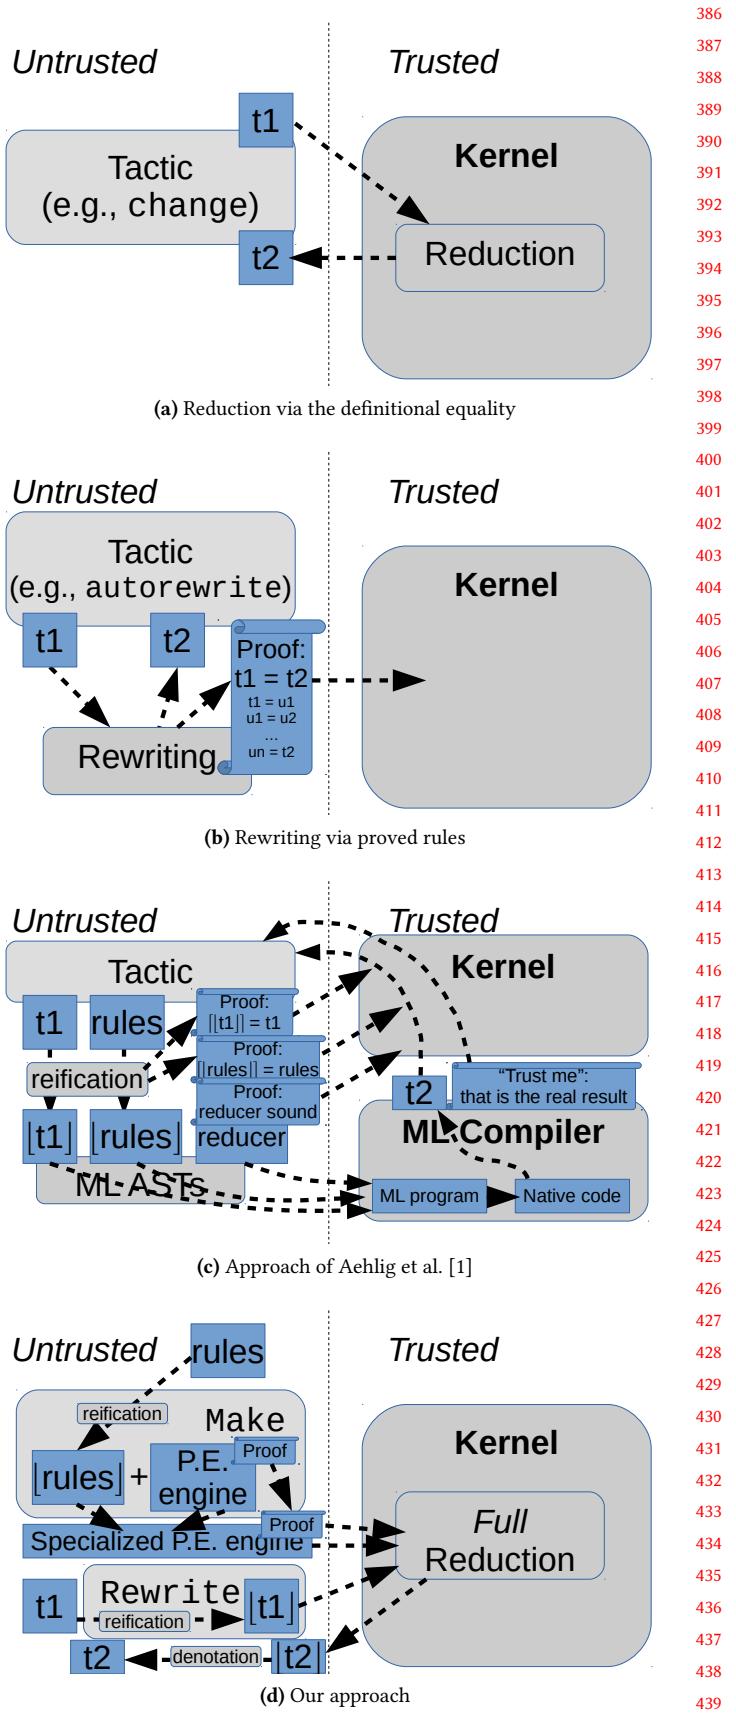


Figure 1. Different approaches to reduction and rewriting

441 is then passed back over to the tactic, in terms of an axiomatic  
 442 assertion that the ML semantics really yields that answer.

443 As Aehlig et al. [1] argue, their use of external compilation  
 444 and evaluation of ML code adds no real complexity on  
 445 top of that required by the proof assistant – after all, the  
 446 proof assistant itself must be compiled and executed some-  
 447 how. However, the perceived increase of trusted code base  
 448 is not spurious: it is one thing to trust that the toolchain and  
 449 execution environment used by the proof assistant and the  
 450 partial evaluator are well-behaved, and another to rely on  
 451 two descriptions of ML (one deeply embedded in the proof  
 452 assistant and another implied by the compiler) to agree on  
 453 every detail of the semantics. Furthermore, there still is new  
 454 trusted code to translate from the deeply embedded ML sub-  
 455 set into the concrete syntax of the full-scale ML language.  
 456 The vast majority of proof-assistant developments today rely  
 457 on no such embeddings with associated mechanized seman-  
 458 tics, so need we really add one to a proof-checking kernel to  
 459 support efficient partial evaluation?

460 Our answer, diagrammed in Figure 1d, shows a different  
 461 way. We still reify terms and rules into a deeply embedded  
 462 language. However, *the reduction engine is implemented di-  
 463 rectly in the logic*, rather than as a deeply embedded syntax  
 464 tree of an ML program. As a result, the kernel’s own reduc-  
 465 tion engine is prepared to execute our reduction engine for  
 466 us – using an operation that would be included in a type-  
 467 theoretic proof assistant in any case, with no special support  
 468 for a language deep embedding. We also stage the process  
 469 for performance reasons. First, the Make command creates  
 470 a rewriter out of a list of rewrite rules, by specializing a  
 471 generic partial-evaluation engine, which has a generic proof  
 472 that applies to any set of proved rewrite rules. We perform  
 473 partial evaluation on the specialized partial evaluator, using  
 474 Coq’s normal reduction mechanisms, under the theory that  
 475 we can afford to pay performance costs at this stage because  
 476 we only need to create new rewriters relatively infrequently.  
 477 Then individual rewritings involve reifying terms, asking  
 478 the kernel to execute the specialized evaluator on them, and  
 479 simplifying an application of an interpretation function to  
 480 the result (this last step must be done using Coq’s normal  
 481 reduction, and it is the bottleneck for outputs with enormous  
 482 numbers of nested binders as discussed in section 5.1).

## 484 2.1 Our Approach in Nine Steps

485 Here is a bit more detail on the steps that go into applying our  
 486 Coq plugin, many of which we expand on in the following  
 487 sections. In order to build a precomputed rewriter with the  
 488 Make command, the following actions are performed:

- 490 1. The given lemma statements are scraped for which  
   491   named functions and types the rewriter package will  
   492   support.
- 493 2. Inductive types enumerating all available primitive  
   494   types and functions are emitted.

495 3. Tactics generate all of the necessary definitions and  
 496 prove all of the necessary lemmas for dealing with this  
 497 particular set of inductive codes. Definitions include  
 498 operations like Boolean equality on type codes and  
 499 lemmas like “all representable primitive types have  
 500 decidable equality.”

- 501 4. The statements of rewrite rules are reified, and we  
   502 prove soundness and syntactic-well-formedness lemm-  
   503 as about each of them. Each instance of the former  
   504 involves wrapping the user-provided proof with the  
   505 right adapter to apply to the reified version.
- 506 5. The definitions needed to perform reification and rewrit-  
   507 ing and the lemmas needed to prove correctness are  
   508 assembled into a single package that can be passed by  
   509 name to the rewriting tactic.

510 When we want to rewrite with a rewriter package in a  
 511 goal, the following steps are performed:

- 512 1. We rearrange the goal into a single logical formula:  
   513   all free-variable quantification in the proof context is  
   514   replaced by changing the equality goal into an equality  
   515   between two functions (taking the free variables as  
   516   inputs).
- 517 2. We reify the side of the goal we want to simplify, using  
   518   the inductive codes in the specified package. That side  
   519   of the goal is then replaced with a call to a denotation  
   520   function on the reified version.
- 521 3. We use a theorem stating that rewriting preserves  
   522   denotations of well-formed terms to replace the de-  
   523   notation subterm with the denotation of the rewriter  
   524   applied to the same reified term. We use Coq’s built-in  
   525   full reduction (`vm_compute`) to reduce the application  
   526   of the rewriter to the reified term.
- 527 4. Finally, we run `cbv` (a standard call-by-value reducer)  
   528   to simplify away the invocation of the denotation func-  
   529   tion on the concrete syntax tree from rewriting.

## 533 3 The Structure of a Rewriter

534 We now simultaneously review the approach of Aehlig et al.  
 535 [1] and introduce some notable differences in our own ap-  
 536 proach, noting similarities to the reflective rewriter of Malecha  
 537 and Bengtson [18] where applicable.

538 First, let us describe the language of terms we support  
 539 rewriting in. Note that, while we support rewriting in full-  
 540 scale Coq proofs, where the metalanguage is dependently  
 541 typed, the object language of our rewriter is nearly simply  
 542 typed, with limited support for calling polymorphic func-  
 543 tions. However, we still support identifiers whose definitions  
 544 use dependent types, since our reducer does not need to look  
 545 into definitions.

$$e ::= \text{App } e_1 \ e_2 \mid \text{Let } v = e_1 \ \text{In } e_2 \\ \mid \text{Abs } (\lambda v. e) \mid \text{Var } v \mid \text{Ident } i$$

546  
 547  
 548  
 549  
 550

551 The Ident case is for identifiers, which are described by an  
 552 enumeration specific to a use of our library. For example, the  
 553 identifiers might be codes for  $+$ ,  $\cdot$ , and literal constants. We  
 554 write  $\llbracket e \rrbracket$  for a standard denotational semantics.  
 555

### 556 3.1 Pattern-Matching Compilation and Evaluation

558 Aehlig et al. [1] feed a specific set of user-provided rewrite  
 559 rules to their engine by generating code for an ML function,  
 560 which takes in deeply embedded term syntax (actually  
 561 *doubly* deeply embedded, within the syntax of the deeply em-  
 562 bedded ML!) and uses ML pattern matching to decide which  
 563 rule to apply at the top level. Thus, they delegate efficient  
 564 implementation of pattern matching to the underlying ML  
 565 implementation. As we instead build our rewriter in Coq's  
 566 logic, we have no such option to defer to ML. Indeed, Coq's  
 567 logic only includes primitive pattern-matching constructs to  
 568 match one constructor at a time.

569 We could follow a naive strategy of repeatedly matching  
 570 each subterm against a pattern for every rewrite rule, as in  
 571 the rewriter of Malecha and Bengtson [18], but in that case  
 572 we do a lot of duplicate work when rewrite rules use overlap-  
 573 ping function symbols. Instead, we adopted the approach of  
 574 Maranget [20], who describes compilation of pattern matches  
 575 in OCaml to decision trees that eliminate needless repeated  
 576 work (for example, decomposing an expression into  $x + y + z$   
 577 only once even if two different rules match on that pattern).  
 578 We have not yet implemented any of the optimizations de-  
 579 scribed therein for finding *minimal* decision trees.

580 There are three steps to turn a set of rewrite rules into a  
 581 functional program that takes in an expression and reduces  
 582 according to the rules. The first step is pattern-matching com-  
 583 pilation: we must compile the lefthand sides of the rewrite  
 584 rules to a decision tree that describes how and in what order  
 585 to decompose the expression, as well as describing which  
 586 rewrite rules to try at which steps of decomposition. Because  
 587 the decision tree is merely a decomposition hint, we require  
 588 no proofs about it to ensure soundness of our rewriter. The  
 589 second step is decision-tree evaluation, during which we  
 590 decompose the expression as per the decision tree, select-  
 591 ing which rewrite rules to attempt. The only correctness  
 592 lemma needed for this stage is that any result it returns is  
 593 equivalent to picking some rewrite rule and rewriting with  
 594 it. The third and final step is to actually rewrite with the  
 595 chosen rule. Here the correctness condition is that we must  
 596 not change the semantics of the expression. Said another  
 597 way, any rewrite-rule replacement expression must match  
 598 the semantics of the rewrite-rule pattern.

599 While pattern matching begins with comparing one pat-  
 600 tern against one expression, Maranget's approach works  
 601 with intermediate goals that check multiple patterns against  
 602 multiple expressions. A decision tree describes how to match  
 603 a vector (or list) of patterns against a vector of expressions.  
 604 It is built from these constructors:

- TryLeaf  $k$  onfailure: Try the  $k^{\text{th}}$  rewrite rule; if it  
 606 fails, keep going with onfailure.
- Failure: Abort; nothing left to try.
- Switch  $\text{icases}$   $\text{app\_case}$  default: With the first  
 609 element of the vector, match on its kind; if it is an  
 610 identifier matching something in  $\text{icases}$ , remove the  
 611 first element of the vector and run that decision tree; if  
 612 it is an application and  $\text{app\_case}$  is not None, try the  
 613  $\text{app\_case}$  decision tree, replacing the first element of  
 614 each vector with the two elements of the function and  
 615 the argument it is applied to; otherwise, do not modify  
 616 the vectors and use the  $\text{default}$  decision tree.
- Swap  $i$  cont: Swap the first element of the vector  
 618 with the  $i^{\text{th}}$  element (0-indexed) and keep going with  
 619 cont.

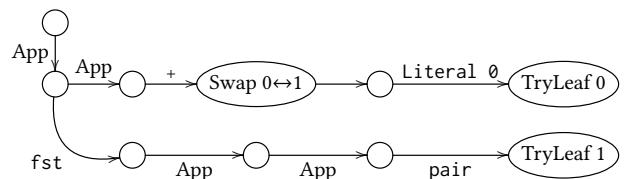
621 Consider the encoding of two simple example rewrite  
 622 rules, where we follow Coq's  $\mathcal{L}_{\text{tac}}$  language in prefacing  
 623 pattern variables with question marks.

$$\begin{aligned} ?n + 0 &\rightarrow n \\ \text{fst}_{\mathbb{Z}, \mathbb{Z}}(?x, ?y) &\rightarrow x \end{aligned}$$

624 We embed them in an AST type for patterns, which largely  
 625 follows our ASTs for expressions.

0. App (App (Ident  $+$ ) Wildcard) (Ident (Literal 0))
1. App (Ident fst) (App (App (Ident pair) Wildcard) Wildcard)

633 The decision tree produced is



641 where every non-swap node implicitly has a “default” case  
 642 arrow to Failure.

643 We implement, in Coq's logic, an evaluator for these trees  
 644 against terms. Note that we use Coq's normal partial eval-  
 645 uation to turn our general decision-tree evaluator into a  
 646 specialized matcher to get reasonable efficiency. Although  
 647 this partial evaluation of our partial evaluator is subject to  
 648 the same performance challenges we highlighted in the in-  
 649 troduction, it only has to be done once for each set of rewrite  
 650 rules, and we are targeting cases where the time of per-goal  
 651 reduction dominates this time of meta-compilation.

652 For our running example of two rules, specializing gives us  
 653 this match expression.

```

match e with
| App f y => match f with
| Ident fst => match y with
| App (App (Ident pair) x) y => x
| _ => e end
| App (Ident +) x => match y with
  
```

```

661 | Ident (Literal 0) => x | _ => e end
662 | _ => e end | _ => e end.
663

```

### 3.2 Adding Higher-Order Features

Fast rewriting at the top level of a term is the key ingredient for supporting customized algebraic simplification. However, not only do we want to rewrite throughout the structure of a term, but we also want to integrate with simplification of higher-order terms, in a way where we can prove to Coq that our syntax-simplification function always terminates. Normalization by evaluation (NbE) [4] is an elegant technique for adding the latter aspect, in a way where we avoid needing to implement our own  $\lambda$ -term reducer or prove it terminating.

To orient expectations: we would like to enable the following reduction

$$(\lambda f\ x\ y.\ f\ x\ y)\ (+)\ z\ 0 \rightsquigarrow z$$

using the rewrite rule

$$?n + 0 \rightarrow n$$

Aehlig et al. [1] also use NbE, and we begin by reviewing its most classic variant, for performing full  $\beta$ -reduction in a simply typed term in a guaranteed-terminating way. The simply typed  $\lambda$ -calculus syntax we use is:

$$t ::= t \rightarrow t \mid b \quad e ::= \lambda v. e \mid e\ e \mid v \mid c$$

with  $v$  for variables,  $c$  for constants, and  $b$  for base types.

We can now define normalization by evaluation. First, we choose a “semantic” representation for each syntactic type, which serves as the result type of an intermediate interpreter.

$$\text{NbE}_t(t_1 \rightarrow t_2) = \text{NbE}_t(t_1) \rightarrow \text{NbE}_t(t_2)$$

$$\text{NbE}_t(b) = \text{expr}(b)$$

Function types are handled as in a simple denotational semantics, while base types receive the perhaps-counterintuitive treatment that the result of “executing” one is a syntactic expression of the same type. We write  $\text{expr}(b)$  for the metalanguage type of object-language syntax trees of type  $b$ , relying on a dependent type family  $\text{expr}$ .

Now the core of NbE, shown in Figure 2, is a pair of dual functions  $\text{reify}$  and  $\text{reflect}$ , for converting back and forth between syntax and semantics of the object language, defined by primitive recursion on type syntax. We split out analysis of term syntax in a separate function  $\text{reduce}$ , defined by primitive recursion on term syntax, when usually this functionality would be mixed in with  $\text{reflect}$ . The reason for this choice will become clear when we extend NbE to handle our full problem domain.

We write  $v$  for object-language variables and  $x$  for metalanguage (Coq) variables, and we overload  $\lambda$  notation using the metavariable kind to signal whether we are building a host  $\lambda$  or a  $\lambda$  syntax tree for the embedded language. The crucial first clause for  $\text{reduce}$  replaces object-language variable

reify <sub>t</sub>	: NbE <sub>t</sub> (t) → expr(t)	716
reify <sub>t<sub>1</sub>→t<sub>2</sub></sub>	(f) = $\lambda v.$ reify <sub>t<sub>2</sub></sub> (f(reflect <sub>t<sub>1</sub></sub> (v)))	717
reify <sub>b</sub>	(f) = f	718
reflect <sub>t</sub>	: expr(t) → NbE <sub>t</sub> (t)	719
reflect <sub>t<sub>1</sub>→t<sub>2</sub></sub>	(e) = $\lambda x.$ reflect <sub>t<sub>2</sub></sub> (e(reify <sub>t<sub>1</sub></sub> (x)))	720
reflect <sub>b</sub>	(e) = e	721
reduce	: expr(t) → NbE <sub>t</sub> (t)	722
reduce( $\lambda v.$ e)	= $\lambda x.$ reduce([x/v]e)	723
reduce(e <sub>1</sub> e <sub>2</sub> )	= (reduce(e <sub>1</sub> )) (reduce(e <sub>2</sub> ))	724
reduce(x)	= x	725
reduce(c)	= reflect(c)	726
NbE	: expr(t) → expr(t)	727
NbE(e)	= reify(reduce(e))	728

**Figure 2.** Implementation of normalization by evaluation

$v$  with fresh metalanguage variable  $x$ , and then we are somehow tracking that all free variables in an argument to  $\text{reduce}$  must have been replaced with metalanguage variables by the time we reach them. We reveal in subsection 4.1 the encoding decisions that make all the above legitimate, but first let us see how to integrate use of the rewriting operation from the previous section. To fuse NbE with rewriting, we only modify the constant case of  $\text{reduce}$ . First, we bind our specialized decision-tree engine under the name  $\text{rewrite-head}$ . Recall that this function only tries to apply rewrite rules at the top level of its input.

In the constant case, we still reflect the constant, but underneath the binders introduced by full  $\eta$ -expansion, we perform one instance of rewriting. In other words, we change this one function-definition clause:

$$\text{reflect}_b(e) = \text{rewrite-head}(e)$$

It is important to note that a constant of function type will be  $\eta$ -expanded only once for each syntactic occurrence in the starting term, though the expanded function is effectively a thunk, waiting to perform rewriting again each time it is called. From first principles, it is not clear why such a strategy terminates on all possible input terms, though we work up to convincing Coq of that fact.

The details so far are essentially the same as in the approach of Aehlig et al. [1]. Recall that their rewriter was implemented in a deeply embedded ML, while ours is implemented in Coq’s logic, which enforces termination of all functions. Aehlig et al. did not prove termination, which indeed does not hold for their rewriter in general, which works with untyped terms, not to mention the possibility of

rule-specific ML functions that diverge themselves. In contrast, we need to convince Coq up-front that our interleaved  $\lambda$ -term normalization and algebraic simplification always terminate. Additionally, we need to prove that our rewriter preserves denotations of terms, which can easily devolve into tedious binder bookkeeping, depending on encoding.

The next section introduces the techniques we use to avoid explicit termination proof or binder bookkeeping, in the context of a more general analysis of scaling challenges.

## 4 Scaling Challenges

Aehlig et al. [1] only evaluated their implementation against closed programs. What happens when we try to apply the approach to partial-evaluation problems that should generate thousands of lines of low-level code?

### 4.1 Variable Environments Will Be Large

We should think carefully about representation of ASTs, since many primitive operations on variables will run in the course of a single partial evaluation. For instance, Aehlig et al. [1] reported a significant performance improvement changing variable nodes from using strings to using de Bruijn indices [9]. However, de Bruijn indices and other first-order representations remain painful to work with. We often need to fix up indices in a term being substituted in a new context. Even looking up a variable in an environment tends to incur linear time overhead, thanks to traversal of a list. Perhaps we can do better with some kind of balanced-tree data structure, but there is a fundamental performance gap versus the arrays that can be used in imperative implementations. Unfortunately, it is difficult to integrate arrays soundly in a logic. Also, even ignoring performance overheads, tedious binder bookkeeping complicates proofs.

Our strategy is to use a variable encoding that pushes all first-order bookkeeping off on Coq's kernel, which is itself performance-tuned with some crucial pieces of imperative code. Parametric higher-order abstract syntax (PHOAS) [8] is a dependently typed encoding of syntax where binders are managed by the enclosing type system. It allows for relatively easy implementation and proof for NbE, so we adopted it for our framework.

Here is the actual inductive definition of term syntax for our object language, PHOAS-style. The characteristic oddity is that the core syntax type `expr` is parameterized on a dependent type family for representing variables. However, the final representation type `Expr` uses first-class polymorphism over choices of variable type, bootstrapping on the metalanguage's parametricity to ensure that a syntax tree is agnostic to variable type.

```

Inductive type := arrow (s d : type)
| base (b : base_type).
Infix " $\rightarrow$ " := arrow.
Inductive expr (var : type  $\rightarrow$  Type)
: type  $\rightarrow$  Type :=

```

```

| Var {t} (v : var t) : expr var t
| Abs {s d} (f : var s  $\rightarrow$  expr var d)
: expr var (s  $\rightarrow$  d)
| App {s d} (f : expr var (s  $\rightarrow$  d))
(x : expr var s) : expr var d
| Const {t} (c : const t) : expr var t
Definition Expr (t : type) : Type :=
forall var, expr var t.

```

A good example of encoding adequacy is assigning a simple denotational semantics. First, a simple recursive function assigns meanings to types.

```

Fixpoint denoteT (t : type) : Type
:= match t with
| arrow s d => denoteT s  $\rightarrow$  denoteT d
| base b     => denote_base_type b
end.

```

Next we see the convenience of being able to *use* an expression by choosing how it should represent variables. Specifically, it is natural to choose *the type-denotation function itself* as the variable representation. Especially note how this choice makes rigorous the convention we followed in the prior section, where a recursive function enforces that values have always been substituted for variables early enough.

```

Fixpoint denoteE {t} (e : expr denoteT t) : denoteT t
:= match e with
| Var v      => v
| Abs f       =>  $\lambda$  x, denoteE (f x)
| App f x    => (denoteE f) (denoteE x)
| Ident c    => denoteI c
end.

```

```

Definition DenoteE {t} (E : Expr t) : denoteT t
:= denoteE (E denoteT).

```

It is now easy to follow the same script in making our rewriting-enabled NbE fully formal. Note especially the first clause of `reduce`, where we avoid variable substitution precisely because we have chosen to represent variables with normalized semantic values. The subtlety there is that base-type semantic values are themselves expression syntax trees, which depend on a nested choice of variable representation, which we retain as a parameter throughout these recursive functions. The final definition  $\lambda$ -quantifies over that choice.

```

Fixpoint nbeT var (t : type) : Type
:= match t with
| arrow s d => nbeT var s  $\rightarrow$  nbeT var d
| base b     => expr var b
end.
Fixpoint reify {var t} : nbeT var t  $\rightarrow$  expr var t
:= match t with
| arrow s d =>  $\lambda$  f,
Abs ( $\lambda$  x, reify (f (reflect (Var x))))
| base b     =>  $\lambda$  e, e
end
with reflect {var t} : expr var t  $\rightarrow$  nbeT var t
:= match t with
| arrow s d =>  $\lambda$  e,

```

```

881       $\lambda x, \text{reflect} (\text{App } e (\text{reify } x))$ 
882      | base b    => rewrite_head
883      end.
884 Fixpoint reduce {var t}
885   (e : expr (nbeT var) t) : nbeT var t
886   := match e with
887   | Abs e      =>  $\lambda x, \text{reduce} (e (\text{Var } x))$ 
888   | App e1 e2 => (reduce e1) (reduce e2)
889   | Var x     => x
890   | Ident c   => reflect (Ident c)
891   end.
892 Definition Rewrite {t} (E : Expr t) : Expr t
893   :=  $\lambda \text{var}, \text{reify} (\text{reduce} (E (\text{nbeT var } t)))$ .
894

```

One subtlety hidden above in implicit arguments is in the final clause of `reduce`, where the two applications of the `Ident` constructor use different variable representations. With all those details hashed out, we can prove a pleasingly simple correctness theorem, with a lemma for each main definition, with inductive structure mirroring recursive structure of the definition, also appealing to correctness of last section's pattern-compilation operations.

$$\forall t, E : \text{Expr } t. \llbracket \text{Rewrite}(E) \rrbracket = \llbracket E \rrbracket$$

Even before getting to the correctness theorem, we needed to convince Coq that the function terminates. While for Aehlig et al. [1], a termination proof would have been a whole separate enterprise, it turns out that PHOAS and NbE line up so well that Coq accepts the above code with no additional termination proof. As a result, the Coq kernel is ready to run our `Rewrite` procedure during checking.

To understand how we now apply the soundness theorem in a tactic, it is important to note that the Coq kernel's built-in reduction strategies have, to an extent, been tuned to work well to show equivalence between a simple denotational-semantics application and the semantic value it produces, while it is rather difficult to code up one reduction strategy that works well for all partial-evaluation tasks. Therefore, we should restrict ourselves to (1) running full reduction in the style of functional-language interpreters and (2) running normal reduction on "known-good" goals like correctness of evaluation of a denotational semantics on a concrete input.

Operationally, then, we apply our tactic in a goal containing a term  $e$  that we want to partially evaluate. In standard proof-by-reflection style, we *reify*  $e$  into some  $E$  where  $\llbracket E \rrbracket = e$ , replacing  $e$  accordingly, asking Coq's kernel to validate the equivalence via standard reduction. Now we use the `Rewrite` correctness theorem to replace  $\llbracket E \rrbracket$  with  $\llbracket \text{Rewrite}(E) \rrbracket$ . Next we may ask the Coq kernel to simplify `Rewrite(E)` by *full reduction via compilation to native code*, since we carefully designed `Rewrite(E)` and its dependencies to produce closed syntax trees. Finally, where  $E'$  is the result of that reduction, we simplify  $\llbracket E' \rrbracket$  with standard reduction, producing a normal-looking Coq term.

## 4.2 Subterm Sharing is Crucial

For some large-scale partial-evaluation problems, it is important to represent output programs with sharing of common subterms. Redundantly inlining shared subterms can lead to exponential increase in space requirements. Consider the Fiat Cryptography [11] example of generating a 64-bit implementation of field arithmetic for the P-256 elliptic curve. The library has been converted manually to continuation-passing style, allowing proper generation of `let` binders, whose variables are often mentioned multiple times. We ran their code generator (actually just a subset of its functionality, but optimized by us a bit further, as explained in subsection 5.2) on the P-256 example and found it took about 15 seconds to finish. Then we modified reduction to inline `let` binders instead of preserving them, at which point the reduction job terminated with an out-of-memory error, on a machine with 64 GB of RAM. (The successful run uses under 2 GB.)

We see a tension here between performance and nice-ness of library implementation. The Fiat Cryptography authors found it necessary to CPS-convert their code to coax Coq into adequate reduction performance. Then all of their correctness theorems were complicated by reasoning about continuations. It feels like a slippery slope on the path to implementing a domain-specific compiler, rather than taking advantage of the pleasing simplicity of partial evaluation on natural functional programs. Our reduction engine takes shared-subterm preservation seriously while applying to libraries in direct style.

Our approach is `let`-lifting: we lift `lets` to top level, so that applications of functions to `lets` are available for rewriting. For example, we can perform the rewriting

$$\begin{aligned} & \text{map } (\lambda x. y + x) (\text{let } z := e \text{ in } [0; 1; 2; z; z + 1]) \\ & \rightsquigarrow \text{let } z := e \text{ in } [y; y + 1; y + 2; y + z; y + (z + 1)] \end{aligned}$$

using the rules

$$\begin{aligned} \text{map } ?f [] &\rightarrow [] & ?n + 0 &\rightarrow n \\ \text{map } ?f (?x ::?xs) &\rightarrow f x :: \text{map } f xs \end{aligned}$$

Our approach is to define a telescope-style type family called `UnderLets`:

```

Inductive UnderLets {var} (T : Type) :=
| Base (v : T)
| UnderLet {A}(e : @expr var A)(f : var A -> UnderLets T).

```

A value of type `UnderLets T` is a series of `let` binders (where each expression  $e$  may mention earlier-bound variables) ending in a value of type  $T$ . It is easy to build various "smart constructors" working with this type, for instance to construct a function application by lifting the `lets` of both function and argument to a common top level.

Such constructors are used to implement an NbE strategy that outputs `UnderLets` telescopes. Recall that the NbE type interpretation mapped base types to expression syntax trees.

```

991 We now parameterize that type interpretation by a Boolean
992 declaring whether we want to introduce telescopes.
993 Fixpoint nbeT' {var} (with_lets : bool) (t : type)
994   := match t with
995     | base t => if with_lets
996       then @UnderLets var (@expr var t)
997       else @expr var t
998     | arrow s d => nbeT' false s -> nbeT' true d
999     end.
1000 Definition nbeT := nbeT' false.
1001 Definition nbeT_with_lets := nbeT' true.

```

There are cases where naive preservation of let binders leads to suboptimal performance, so we include some heuristics. For instance, when the expression being bound is a constant, we always inline. When the expression being bound is a series of list “cons” operations, we introduce a name for each individual list element, since such a list might be traversed multiple times in different ways.

### 4.3 Rules Need Side Conditions

Many useful algebraic simplifications require side conditions. One simple case is supporting *nonlinear* patterns, where a pattern variable appears multiple times. We can encode nonlinearity on top of linear patterns via side conditions.

$$\text{?}n_1 + \text{?}m - \text{?}n_2 \rightarrow m \text{ if } n_1 = n_2$$

The trouble is how to support predictable solving of side conditions during partial evaluation, where we may be rewriting in open terms. We decided to sidestep this problem by allowing side conditions only as executable Boolean functions, to be applied only to variables that are confirmed as *compile-time constants*, unlike Malecha and Bengtson [18] who support general unification variables. We added a variant of pattern variable that only matches constants. Semantically, this variable style has no additional meaning, and in fact we implement it as a special identity function that should be called in the right places within Coq lemma statements. Rather, use of this identity function triggers the right behavior in our tactic code that reifies lemma statements. We introduce a notation where a prefixed apostrophe signals a call to the “constants only” function.

Our reification inspects the hypotheses of lemma statements, using type classes to find decidable realizations of the predicates that are used, synthesizing one Boolean expression of our deeply embedded term language, standing for a decision procedure for the hypotheses. The Make command fails if any such expression contains pattern variables not marked as constants. Therefore, matching of rules can safely run side conditions, knowing that Coq’s full-reduction engine can determine their truth efficiently.

### 4.4 Side Conditions Need Abstract Interpretation

With our limitation that side conditions are decided by executable Boolean procedures, we cannot yet handle directly

some of the rewrites needed for realistic partial evaluation. For instance, Fiat Cryptography reduces high-level functional to low-level code that only uses integer types available on the target hardware. The starting library code works with infinite-precision integers, while the generated low-level code should be careful to avoid unintended integer overflow. As a result, the setup may be too naive for our running example rule  $?n + 0 \rightarrow n$ . When we get to reducing fixed-precision-integer terms, we must be legalistic:

$$\text{add\_with\_carry}_{64}(\text{?}n, 0) \rightarrow (0, n) \text{ if } 0 \leq n < 2^{64}$$

We developed a design pattern to handle this kind of rule.

First, we introduce a family of functions  $\text{clip}_{l,u}$ , each of which forces its integer argument to respect lower bound  $l$  and upper bound  $u$ . Partial evaluation is proved with respect to unknown realizations of these functions, only requiring that  $\text{clip}_{l,u}(n) = n$  when  $l \leq n < u$ . Now, before we begin partial evaluation, we can run a verified abstract interpreter to find conservative bounds for each program variable. When bounds  $l$  and  $u$  are found for variable  $x$ , it is sound to replace  $x$  with  $\text{clip}_{l,u}(x)$ . Therefore, at the end of this phase, we assume all variable occurrences have been rewritten in this manner to record their proved bounds.

Second, we proceed with our example rule refactored:

$$\text{add\_with\_carry}_{64}(\text{clip}_{?l,?u}(\text{?}n), 0) \rightarrow (0, \text{clip}_{l,u}(n))$$

if  $u < 2^{64}$

If the abstract interpreter did its job, then all lower and upper bounds are constants, and we can execute side conditions straightforwardly during pattern matching.

## 5 Evaluation

Our implementation, attached to this submission as an anonymized supplement with a roadmap in Appendix D, includes a mix of Coq code for the proved core of rewriting, tactic code for setting up proper use of that core, and OCaml plugin code for the manipulations beyond the current capabilities of the tactic language. We report here on experiments to isolate performance benefits for rewriting under binders and reducing higher-order structure.

### 5.1 Microbenchmarks

We start with microbenchmarks focusing attention on particular aspects of reduction and rewriting, with Appendix A going into more detail.

#### 5.1.1 Rewriting Under Binders

Consider

let	$v_1 := v_0 + v_0 + 0$	in
:		
let	$v_n := v_{n-1} + v_{n-1} + 0$	in
	$v_n + v_n + 0$	

We want to remove all of the  $+ 0$ s. We can start from this expression directly, in which case reification alone takes as

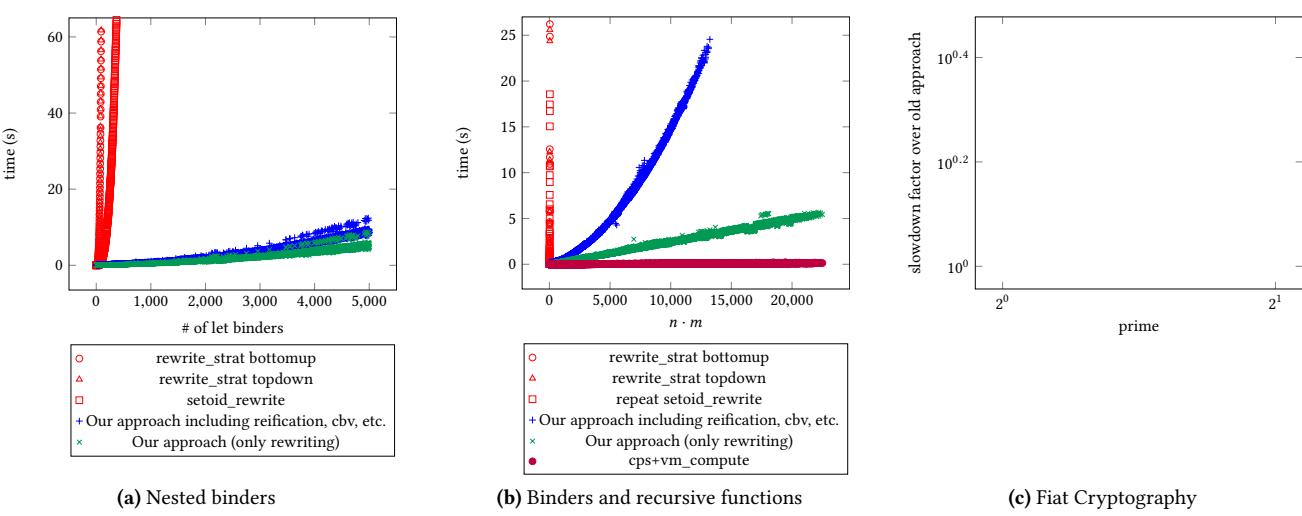


Figure 3. Timing of different partial-evaluation implementations

much time as `setoid_rewrite`. As the reification method was not especially optimized, and there exist fast reification methods [12], we instead start from a call to a recursive function that generates such a sequence of `let` bindings.

Figure 3a shows the results. The comparison points are Coq’s `setoid_rewrite` and `rewrite_strat`. The former performs one rewrite at a time, taking minimal advantage of commonalities across them and thus generating quite large, redundant proof terms. The latter makes top-down or bottom-up passes with combined generation of proof terms. For our own approach, we list both the total time and the time taken for core execution of a verified rewrite engine, without counting reification (converting goals to ASTs) or its inverse (interpreting results back to normal-looking goals).

The comparison here is very favorable for our approach. The competing tactics spike upward toward timeouts at just a few hundred generated binders, while our engine is only taking about 10 seconds for examples with 5,000 nested binders.

As detailed in subsection A.2, we ran a variant of this experiment with inlining of `lets`, forcing terms to grow quite large. Specifically, we generate  $n$  nested `lets`, each repeatedly adding a designated free variable into a sum,  $m$  times. Holding  $m$  fixed at a small value and letting  $n$  scale, we continue dominating the methods described above, though Coq’s `rewrite!` tactic (to rewrite with one lemma many times) does better for  $m < 2$ . Holding  $n$  fixed and letting  $m$  scale, all other approaches quickly spike upward to timeouts, while ours holds steady even for  $m = 1000$ .

### 5.1.2 Binders and Recursive Functions

The next experiment uses the following example.

$$\text{map\_dbl}(\ell) = \begin{cases} [] & \text{if } \ell = [] \\ \text{let } y := h + h \text{ in } y :: \text{map\_dbl}(t) & \text{if } \ell = h :: t \end{cases}$$

$$\text{make}(n, m, v) = \begin{cases} [v, \underbrace{\dots, v}_n] & \text{if } m = 0 \\ \text{map\_dbl}(\text{make}(n, m - 1, v)) & \text{if } m > 0 \end{cases}$$

`examplen, m = ∀v, make(n, m, v) = []`

Note that the `let ... in ...` binding blocks further reduction of `map_dbl`, which we iterate  $m$  times, and so we need to take care to preserve sharing when reducing here.

Figure 3b compares performance between our approach, `repeat_setoid_rewrite`, and two variants of `rewrite_strat`. Additionally, we consider another option, which was adopted by Fiat Cryptography at a larger scale: rewrite our functions to improve reduction behavior. Specifically, both functions are rewritten in continuation-passing style, which makes them harder to read and reason about but allows standard VM-based reduction to achieve good performance. The figure shows that `rewrite_strat` variants are essentially unusable for this example, with `setoid_rewrite` performing only marginally better, while our approach applied to the original, more readable definitions loses ground steadily to VM-based reduction on CPSed code. On the largest terms ( $n \cdot m > 20,000$ ), the gap is 6s vs. 0.1s of compilation time, which should often be acceptable in return for simplified coding and proofs, plus the ability to mix proved rewrite rules with built-in reductions. See subsection A.3 for more on this microbenchmark and subsection A.4 for an even more extreme example of full reduction with a Sieve of Eratosthenes as in the experiments of Ahlig et al. [1] (ours 10s, VM 0.3s).

### 5.2 Macrobenchmark: Fiat Cryptography

Finally, we consider an experiment (described in more detail in Appendix B) replicating the generation of performance-competitive finite-field-arithmetic code for all popular elliptic curves by Erbsen et al. [11]. In all cases, we generate

essentially the same code as they did, so we only measure performance of the code-generation process. We stage partial evaluation with three different reduction engines (i.e., three Make invocations), respectively applying 85, 56, and 44 rewrite rules (with only 2 rules shared across engines), taking total time of about 5 minutes to generate all three engines. These engines support 95 distinct function symbols.

Figure 3c graphs running time of three different partial-evaluation methods for Fiat Cryptography, as the prime modulus of arithmetic scales up. Times are normalized to the performance of the original method, which relied entirely on standard Coq reduction. Actually, in the course of running this experiment, we found a way to improve the old approach for a fairer comparison. It had relied on Coq’s configurable cbv tactic to perform reduction with selected rules of the definitional equality, which the Fiat Cryptography developers had applied to blacklist identifiers that should be left for compile-time execution. By instead hiding those identifiers behind opaque module-signature ascription, we were able to run Coq’s more-optimized virtual-machine-based reducer.

As the figure shows, our approach running partial evaluation inside Coq’s kernel begins with about a 10× performance disadvantage vs. the original method. With log scale on both axes, we see that this disadvantage narrows to become nearly negligible for the largest primes, of around 500 bits. (We used the same set of prime moduli as in the experiments run by Erbsen et al. [11], which were chosen based on searching the archives of an elliptic-curves mailing list for all prime numbers.) It makes sense that execution inside Coq leaves our new approach at a disadvantage, as we are essentially running an interpreter (our normalizer) within an interpreter (Coq’s kernel), while the old approach ran just the latter directly. Also recall that the old approach required rewriting Fiat Cryptography’s library of arithmetic functions in continuation-passing style, enduring this complexity in library correctness proofs, while our new approach applies to a direct-style library. Finally, the old approach included a custom reflection-based arithmetic simplifier for term syntax, run after traditional reduction, whereas now we are able to apply a generic engine that combines both, without requiring more than proving traditional rewrite rules.

The figure also confirms clear performance advantage of running reduction in code extracted to OCaml, which is possible because our plugin produces verified code in Coq’s functional language. By the time we reach middle-of-the-pack prime size around 300 bits, the extracted version is running about 10× as quickly as the baseline.

## 6 Related Work

We have already discussed the work of Aehlig et al. [1], which introduced the basic structure that our engine shares, but which required a substantially larger trusted code base,

did not tackle certain challenges in scaling to large partial-evaluation problems, and did not report any performance experiments in partial evaluation.

We have also mentioned  $\mathcal{R}_{tac}$  [18], which implements an experimental reflective version of `rewrite_strat` supporting arbitrary setoid relations, unification variables, and arbitrary semi-decidable side conditions solvable by other reflective tactics, using de Bruijn indexing to manage binders. We were unfortunately unable to get the rewriter to work with Coq 8.10 and were also not able to determine from the paper how to repurpose the rewriter to handle our benchmarks.

Our implementation builds on fast full reduction in Coq’s kernel, via a virtual machine [13] or compilation to native code [5]. Especially the latter is similar in adopting an NbE style for full reduction, simplifying even under  $\lambda$ s, on top of a more traditional implementation of OCaml that never executes preemptively under  $\lambda$ s. Neither approach unifies support for rewriting with proved rules, and partial evaluation only applies in very limited cases, where functions that should not be evaluated at compile time must have properly opaque definitions that the evaluator will not consult. Neither implementation involved a machine-checked proof suitable to bootstrap on top of reduction support in a kernel providing simpler reduction.

A variety of forms of pragmatic partial evaluation have been demonstrated, with Lightweight Modular Staging [22] in Scala as one of the best-known current examples. A kind of type-based overloading for staging annotations is used to smooth the rough edges in writing code that manipulates syntax trees. The LMS-Verify system [2] can be used for formal verification of generated code after-the-fact. Typically LMS-Verify has been used with relatively shallow properties (though potentially applied to larger and more sophisticated code bases than we tackle), not scaling to the kinds of functional-correctness properties that concern us here, justifying investment in verified partial evaluators.

## 7 Future Work

There are a number of natural extensions to our engine. For instance, we do not yet allow pattern variables marked as “constants only” to apply to container datatypes; we limit the mixing of higher-order and polymorphic types, as well as limiting use of first-class polymorphism; we do not support proving equalities on functions; we only support decidable predicates as rule side conditions, and the predicates may only mention pattern variables restricted to matching constants; we have hardcoded support for a small set of container types and their eliminators; we support rewriting with equality and no other relations (e.g., subset inclusion); and we require decidable equality for all types mentioned in rules. It may be helpful to design an engine that lifts some or all of these limitations, building on the basic structure that we present here.

1266  
1267  
1268  
1269  
1270  
1271  
1272  
1273  
1274  
1275  
1276  
1277  
1278  
1279  
1280  
1281  
1282  
1283  
1284  
1285  
1286  
1287  
1288  
1289  
1290  
1291  
1292  
1293  
1294  
1295  
1296  
1297  
1298  
1299  
1300  
1301  
1302  
1303  
1304  
1305  
1306  
1307  
1308  
1309  
1310  
1311  
1312  
1313  
1314  
1315  
1316  
1317  
1318  
1319  
1320

## References

- [1] Klaus Aehlig, Florian Haftmann, and Tobias Nipkow. 2008. A Compiled Implementation of Normalization by Evaluation. In *Proc. TPHOLs*. 35–46. <http://moscova.inria.fr/~maranget/papers/ml05e-maranget.pdf>
- [2] Nada Amin and Tiark Rompf. 2017. LMS-Verify: Abstraction without Regret for Verified Systems Programming. In *Proc. POPL*.
- [3] Abhishek Anand, Simon Boulier, Cyril Cohen, Matthieu Sozeau, and Nicolas Tabareau. 2018. Towards Certified Meta-Programming with Typed Template-Coq. In *Proc. ITP*.
- [4] U. Berger and H. Schwichtenberg. 1991. An inverse of the evaluation functional for typed lambda -calculus. In [*1991 Proceedings Sixth Annual IEEE Symposium on Logic in Computer Science*. 203–211. <https://doi.org/10.1109/LICS.1991.151645>
- [5] Mathieu Boespflug, Maxime Dénès, and Benjamin Grégoire. 2011. Full Reduction at Full Throttle. In *Proc. CPP*.
- [6] Barry Bond, Chris Hawblitzel, Manos Kapritsos, Rustan Leino, Jay Lorch, Bryan Parno, Ashay Rane, Srinath Setty, and Laure Thompson. 2017. Vale: Verifying High-Performance Cryptographic Assembly Code. In *Proc. USENIX Security*. <http://www.cs.cornell.edu/~laurejt/papers/vale-2017.pdf>
- [7] Samuel Boutin. 1997. Using reflection to build efficient and certified decision procedures. In *Proc. TACS*.
- [8] Adam Chlipala. 2008. Parametric Higher-Order Abstract Syntax for Mechanized Semantics. In *ICFP'08: Proceedings of the 13th ACM SIGPLAN International Conference on Functional Programming*. <http://adam.chlipala.net/papers/PhaosICFP08/>
- [9] Nicolaas Govert De Bruijn. 1972. Lambda calculus notation with nameless dummies, a tool for automatic formula manipulation, with application to the Church-Rosser theorem. In *Indagationes Mathematicae (Proceedings)*, Vol. 75. Elsevier, 381–392.
- [10] Maxime Dénès. 2013. Towards primitive data types for COQ. In *The Coq Workshop 2013* (2013-04-06). [https://coq.inria.fr/files/coq5\\_submission\\_2.pdf](https://coq.inria.fr/files/coq5_submission_2.pdf)
- [11] Andres Erbsen, Jade Philipoom, Jason Gross, Robert Sloan, and Adam Chlipala. 2019. Simple High-Level Code For Cryptographic Arithmetic – With Proofs, Without Compromises. In *IEEE Security & Privacy*. <http://adam.chlipala.net/papers/FiatCryptoSP19/>
- [12] Jason Gross, Andres Erbsen, and Adam Chlipala. 2018. Reification by Parametricity: Fast Setup for Proof by Reflection, in Two Lines of Ltac. In *Proc. ITP*. <http://adam.chlipala.net/papers/ReificationITP18/>
- [13] Benjamin Grégoire and Xavier Leroy. 2002. A compiled implementation of strong reduction. In *Proc. ICFP*.
- [14] Florian Haftmann and Tobias Nipkow. 2007. A Code Generator Framework for Isabelle/HOL. In *Proc. TPHOLs*.
- [15] N.D. Jones, C.K. Gomard, and P. Sestoft. 1993. *Partial Evaluation and Automatic Program Generation*. Prentice Hall International.
- [16] Gerwin Klein, Kevin Elphinstone, Gernot Heiser, June Andronick, David Cock, Philip Derrin, Dhammika Elkaduve, Kai Engelhardt, Rafal Kolanski, Michael Norrish, Thomas Sewell, Harvey Tuch, and Simon Winwood. 2009. seL4: Formal Verification of an OS Kernel. In *Proc. SOSP*. ACM, 207–220.
- [17] Xavier Leroy. 2009. A Formally Verified Compiler Back-end. *J. Autom. Reason.* 43, 4 (Dec. 2009), 363–446. <http://gallium.inria.fr/~xleroy/publi/compcert-backend.pdf>
- [18] Gregory Malecha and Jesper Bengtson. 2016. *Programming Languages and Systems: 25th European Symposium on Programming, ESOP 2016, Held as Part of the European Joint Conferences on Theory and Practice of Software, ETAPS 2016, Eindhoven, The Netherlands, April 2-8, 2016, Proceedings*. Springer Berlin Heidelberg, Berlin, Heidelberg, Chapter Extensible and Efficient Automation Through Reflective Tactics, 532–559. [https://doi.org/10.1007/978-3-662-49498-1\\_21](https://doi.org/10.1007/978-3-662-49498-1_21)
- [19] Gregory Michael Malecha. 2014. *Extensible Proof Engineering in Intensional Type Theory*. Ph.D. Dissertation. Harvard University. <http://gmalecha.github.io/publication/2015/02/01/extensible-proof-engineering-in-intensional-type-theory.html>

## 1431 A Additional Information on 1432 Microbenchmarks

1433 We performed all benchmarks on a 3.5 GHz Core i7 running  
1434 Linux and Coq 8.10.0. We name the subsections here with  
1435 the names that show up in the code supplement.

### 1437 A.1 UnderLetsPlus0

1438 We provide more detail on the “nested binders” microbenchmark of subsubsection 5.1.1 displayed in Figure 3a.  
1439

1440 Recall that we are removing all of the + 0s from  
1441

```
1442     let v1 := v0 + v0 + 0 in
1443         :
1444             let vn := vn-1 + vn-1 + 0 in
1445                 vn + vn + 0
```

1446 The code used to define this microbenchmark is  
1447

```
1448 Definition make_lets_def (n:nat) (v acc : Z) :=
1449   @nat_rect
1450     (fun _ => Z * Z -> Z)
1451     (fun '(v, acc) => acc + acc + v)
1452     (fun _ rec '(v, acc) =>
1453       dlet acc := acc + acc + v in rec (v, acc))
1454   n
1455   (v, acc).
```

1456 We note some details of the rewriting framework that were  
1457 glossed over in the main body of the paper, which are useful  
1458 for using the code: Although the rewriting framework  
1459 does not support dependently typed constants, we can au-  
1460 tomatically preprocess uses of eliminators like `nat_rect`  
1461 and `list_rect` into non-dependent versions. The tactic that  
1462 does this preprocessing is extensible via  $\mathcal{L}_{\text{tac}}$ ’s reassignment  
1463 feature. Since pattern-matching compilation mixed with NbE  
1464 requires knowing how many arguments a constant can be  
1465 applied to, we must internally use a version of the recur-  
1466 sion principle whose type arguments do not contain arrows;  
1467 current preprocessing can handle recursion principles with  
1468 either no arrows or one arrow in the motive. Even though we  
1469 will eventually plug in 0 for  $v$ , we jump through some extra  
1470 hoops to ensure that our rewriter cannot cheat by rewriting  
1471 away the + 0 before reducing the recursion on  $n$ .  
1472

1473 We can reduce this expression in three ways.

#### 1474 A.1.1 Our Rewriter

1475 One lemma is required for rewriting with our rewriter:

1476 Lemma `Z.add_0_r` : `forall z, z + 0 = z`.

1477 Creating the rewriter takes about 12 seconds on the  
1478 machine we used for running the performance experiments:  
1479

```
1480 Make myrew := Rewriter For
1481   (Z.add_0_r, eval_rect nat, eval_rect prod).
```

1482 Recall from subsection 1.1 that `eval_rect` is a definition  
1483 provided by our framework for eagerly evaluating recur-  
1484 sion associated with certain types. It functions by triggering  
1485

1486 typeclass resolution for the lemmas reducing the recursion  
1487 principle associated to the given type. We provide instances  
1488 for `nat`, `prod`, `list`, `option`, and `bool`. Users may add more  
1489 instances if they desire.

### 1490 A.1.2 setoid\_rewrite and rewrite\_strat

1491 To give as many advantages as we can to the preexisting  
1492 work on rewriting, we pre-reduce the recursion on `nats`  
1493 using `cbv` before performing `setoid_rewrite`. (Note that  
1494 `setoid_rewrite` cannot itself perform reduction without  
1495 generating large proof terms, and `rewrite_strat` is not  
1496 currently capable of sequencing reduction with rewriting in-  
1497 ternally due to bugs such as #10923.) Rewriting itself is easy;  
1498 we may use any of `repeat setoid_rewrite Z.add_0_r`,  
1499 `rewrite_strat topdown Z.add_0_r`, or `rewrite_strat`  
1500 `bottomup Z.add_0_r`.

## 1501 A.2 Plus0Tree

1502 This is a version of subsection A.1 without any let binders,  
1503 discussed in subsubsection 5.1.1 but not displayed in Figure 3.  
1504

1505 We use two definitions for this microbenchmark:

```
1506 Definition iter (m : nat) (acc v : Z) :=
1507   @nat_rect
1508     (fun _ => Z * Z -> Z)
1509     (fun acc => acc)
1510     (fun _ rec acc => rec (acc + v))
1511   m
1512   acc.
```

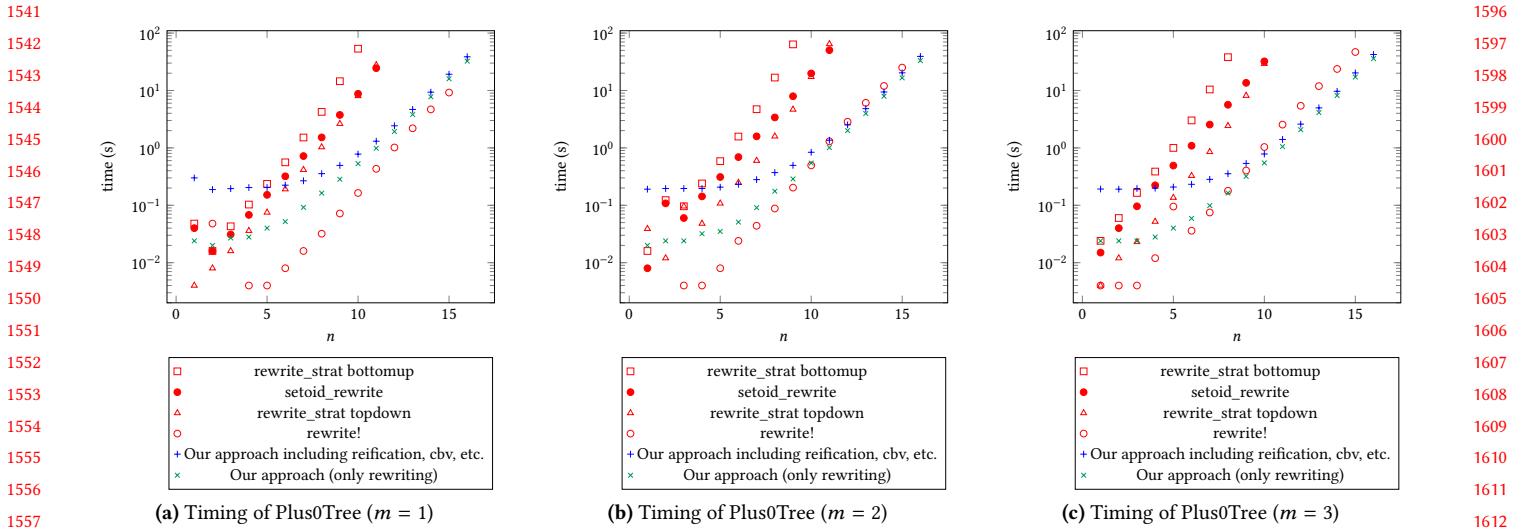
```
1513 Definition make_tree (n m : nat) (v acc : Z) :=
1514   Eval cbv [iter] in
1515   @nat_rect
1516     (fun _ => Z * Z -> Z)
1517     (fun '(v, acc) => iter m (acc + acc) v)
1518     (fun _ rec '(v, acc) =>
1519       iter m (rec (v, acc) + rec (v, acc)) v)
1520   n
1521   (v, acc).
```

1522 We can see from the graphs in Figure 4 and Figure 5 that  
1523 (a) we incur constant overhead over most of the other meth-  
1524 ods which dominates on small examples; (b) when the term  
1525 is quite large and there are few opportunities for rewriting  
1526 relative to the term-size (i.e.,  $m \leq 2$ ), we are worse than  
1527 `rewrite !Z.add_0_r`, but still better than the other meth-  
1528 ods; and (c) when there are many opportunities for rewriting  
1529 relative to the term-size ( $m > 2$ ), we thoroughly dominate  
1530 the other methods.

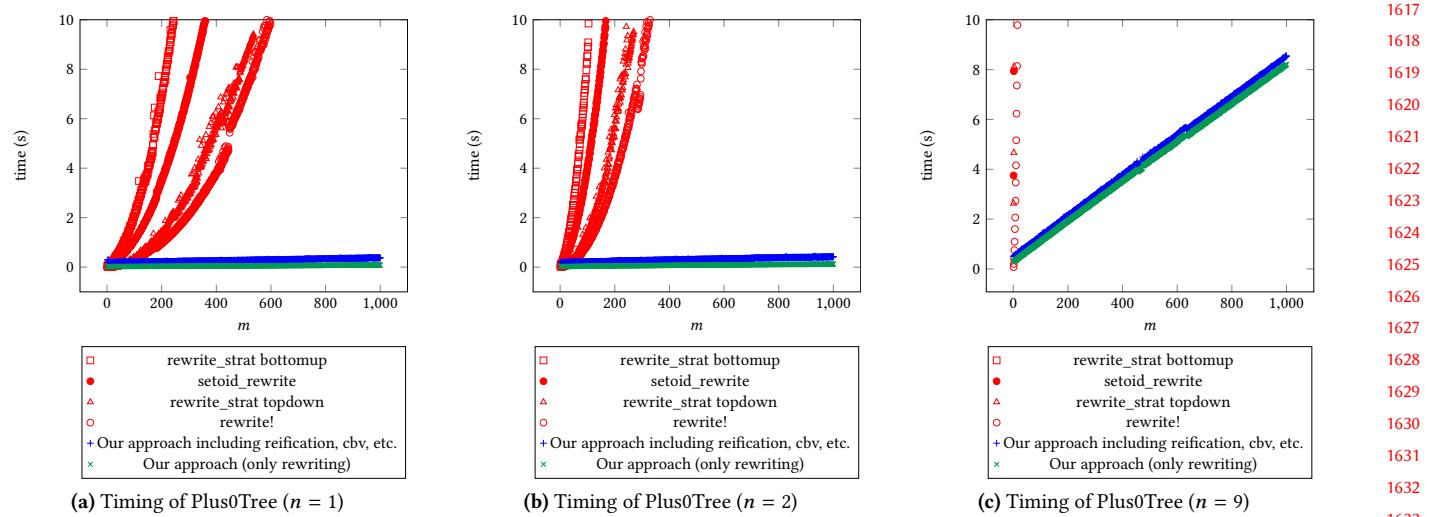
## 1531 A.3 LiftLetsMap

1532 We now discuss in more detail the “binders and recursive  
1533 functions” example from subsubsection 5.1.2.  
1534

1535  
1536  
1537  
1538  
1539  
1540



**Figure 4.** Timing of different partial-evaluation implementations for Plus0Tree for fixed  $m$ . Note that we have a logarithmic time scale, because term size is proportional to  $2^n$ .



**Figure 5.** Timing of different partial-evaluation implementations for Plus0Tree for fixed  $n$  (1, 2, and then we jump to 9)

The expression we want to get out at the end looks like:

```
let v1,1 := v + v in
:
let v1,n := v + v in
let v2,1 := v1,1 + v1,1 in
:
let v2,n := v1,n + v1,n in
:
[vm,1, ..., vm,n]
```

Recall that we make this example with the code

```
Definition map_double (ls : list Z) :=
list_rect
-
[] []
(λ x xs rec, let y := x + x in y :: rec)
ls.

Definition make (n : nat) (m : nat) (v : Z) :=
nat_rect
-
(List.repeat v n)
(λ _ rec, map_double rec)
m.
```

1651 We can perform this rewriting in four ways; see Figure 3b.  
 1652 Note that `rewrite_strat` grows quite quickly, hitting a  
 1653 minute when the total number of rewrites ( $n \cdot m$ ) is in the  
 1654 mid-40s. Our method performs much better, but the fact that  
 1655 we have to perform `cbv` at the end costs us; about 99% of  
 1656 the difference between the full time of our method and just  
 1657 the rewriting is spent in the final `cbv` at the end. This is due  
 1658 to the unfortunate fact that reduction in Coq is quadratic in  
 1659 the number of nested binders present; see Coq bug #11151.  
 1660 Finally, and unsurprisingly, `vm_compute` outperforms us.

### 1661 A.3.1 Our Rewriter

1662 One lemma is required for rewriting with our rewriter:

```
1663 Lemma eval_repeat A x n :
1664   @List.repeat A x ('n)
1665   = ident.eagerly nat_rect _ [
1666     (λ k repeat_k, x :: repeat_k)
1667     ('n).
```

1668 Recall that the apostrophe marker ('') is explained in sub-  
 1669 section 1.1. Recall again from subsection 1.1 that we use  
 1670 `ident.eagerly` to ask the reducer to simplify a case of prim-  
 1671 itive recursion by complete traversal of the designated argu-  
 1672 ment's constructor tree. Our current version only allows a  
 1673 limited, hard-coded set of eliminators with `ident.eagerly`  
 1674 (`nat_rect` on return types with either zero or one arrows,  
 1675 `list_rect` on return types with either zero or one arrows,  
 1676 and `List.nth_default`), but nothing in principle prevents  
 1677 automatic generation of the necessary code.

1678 We construct our rewriter with

```
1679 Make myrew := Rewriter For
1680   (eval_repeat, eval_rect list, eval_rect nat)
1681   (with extra idents (Z.add)).
```

1682 On the machine we used for running all our performance  
 1683 experiments, this command takes about 13 seconds to run.  
 1684 Note that all identifiers which appear in any goal to be rewritten  
 1685 must either appear in the type of one of the rewrite rules  
 1686 or in the tuple passed to `with extra idents`.

1687 Rewriting is relatively simple, now. Simply invoke the  
 1688 tactic `Rewrite_for` `myrew`. We support rewriting on only  
 1689 the left-hand-side and on only the right-hand-side using  
 1690 either the tactic `Rewrite_lhs_for` `myrew` or else the tactic  
 1691 `Rewrite_rhs_for` `myrew`, respectively.

### 1692 A.3.2 `rewrite_strat`

1693 To reduce adequately using `rewrite_strat`, we need the  
 1694 following two lemmas:

```
1695 Lemma lift_let_list_rect T A P N C (v : A) fls
1696   : @list_rect T P N C (Let_In v fls)
1697   = Let_In v (fun v => @list_rect T P N C (fls v)).
1698 Lemma lift_let_cons T A x (v : A) f
1699   : @cons T x (Let_In v f)
1700   = Let_In v (fun v => @cons T x (f v)).
```

1701 Note that `Let_In` is the constant we use for writing `let`  
 1702 ... `in` ... expressions that do not reduce under  $\zeta$ . Through-  
 1703 out most of this paper, anywhere that `let` ... `in` ... ap-  
 1704 pears, we have actually used `Let_In` in the code. It would  
 1705 alternatively be possible to extend the reification preproces-  
 1706 sor to automatically convert `let` ... `in` ... to `Let_In`, but this  
 1707 may cause problems when converting the interpre-  
 1708 tation of the reified term with the pre-reified term, as Coq's  
 1709 conversion does not allow fine-tuning of when to inline or  
 1710 unfold `lets`.

1711 To rewrite, we start with `cbv` [`example make map dbl1`]  
 1712 to expose the underlying term to rewriting. One would  
 1713 hope that one could just add these two hints to a database  
 1714 `db` and then write `rewrite_strat (repeat (eval cbn [list_rect]; try bottomup hints db))`, but unfor-  
 1715 tunately this does not work due to a number of bugs in  
 1716 Coq: #10934, #10923, #4175, #10955, and the potential to  
 1717 hit #10972. Instead, we must put the two lemmas in sepa-  
 1718 rate databases, and then write `repeat (cbn [list_rect];`  
 1719 `(rewrite_strat (try repeat bottomup hints db1));`  
 1720 `(rewrite_strat (try repeat bottomup hints db2)))`. Note that the rewriting with `lift_let_cons` can be done  
 1721 either top-down or bottom-up, but `rewrite_strat` breaks if  
 1722 the rewriting with `lift_let_list_rect` is done top-down.

### 1723 A.3.3 CPS and the VM

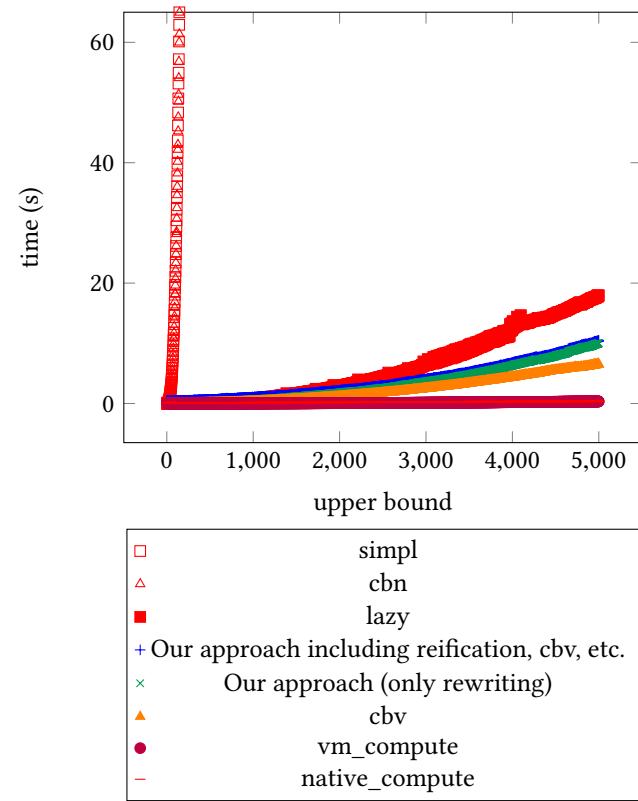
1724 If we want to use Coq's built-in VM reduction without our  
 1725 rewriter, to achieve the prior state-of-the-art performance,  
 1726 we can do so on this example, because it only involves partial  
 1727 reduction and not equational rewriting. However, we must (a)  
 1728 module-opacify the constants which are not to be unfolded,  
 1729 and (b) rewrite all of our code in CPS.

1730 Then we are looking at

$$\begin{aligned} \text{map\_dbl\_cps}(\ell, k) &= \begin{cases} k([]) & \text{if } \ell = [] \\ \text{let } y := h +_{\text{ax}} h \text{ in } & \text{if } \ell = h :: t \\ \text{map\_dbl\_cps}(t, & \\ & (\lambda ys, k(y :: ys))) \end{cases} \\ \text{make\_cps}(n, m, v, k) &= \begin{cases} k(\underbrace{v, \dots, v}_n) & \text{if } m = 0 \\ \text{make\_cps}(n, m - 1, v, & \text{if } m > 0 \\ & (\lambda \ell, \text{map\_dbl\_cps}(\ell, k))) \end{cases} \\ \text{example\_cps}_{n,m} &= \forall v, \text{make\_cps}(n, m, v, \lambda x. x) = [] \end{aligned}$$

1731 Then we can just run `vm_compute`. Note that this strategy,  
 1732 while quite fast, results in a stack overflow when  $n \cdot m$  is  
 1733 larger than approximately  $2.5 \cdot 10^4$ . This is unsurprising, as  
 1734 we are generating quite large terms. Our framework can  
 1735 handle terms of this size but stack-overflows on only slightly  
 1736 larger terms.

1737 1738 1739 1740 1741 1742 1743 1744 1745 1746 1747 1748 1749 1750 1751 1752 1753 1754 1755 1756 1757 1758 1759 1760



**Figure 6.** Timing of different full-evaluation implementations for SieveOfEratosthenes

### A.3.4 Takeaway

From this example, we conclude that `rewrite_strat` is unsuitable for computations involving large terms with many binders, especially in cases where reduction and rewriting need to be interwoven, and that the many bugs in `rewrite_strat` result in confusing gymnastics required for success. The prior state of the art—writing code in CPS—suitably tweaked by using module pacity to allow `vm_compute`, remains the best performer here, though the cost of rewriting everything in CPS may be prohibitive. Our method soundly beats `rewrite_strat`. We are additionally bottlenecked on `cbv`, which is used to unfold the goal post-rewriting and costs about a minute on the largest of terms; see Coq bug #11151 for a discussion on what is wrong with Coq’s reduction here.

### A.4 SieveOfEratosthenes

To benchmark how much overhead we add when we are reducing fully, we compute the Sieve of Eratosthenes, taking inspiration on benchmark choice from Aehlig et al. [1]. We find in Figure 6 that we are slower than `vm_compute`, `native_compute`, and `cbv`, but faster than `lazy`, and of course much faster than `simpl` and `cbn`, which are quite slow.

We define the sieve using `PositiveMap.t` and `list Z`:

```

1761 Definition sieve' (fuel : nat) (max : Z) :=          1816
1762   List.rev                                              1817
1763   (fst                                              1818
1764     (@nat_rect                                         1819
1765       (λ _, list Z (* primes *) *           1820
1766         PositiveSet.t (* composites *) *        1821
1767         positive (* np (next_prime) *) ->    1822
1768         list Z (* primes *) *                  1823
1769         PositiveSet.t (* composites *)           1824
1770         (λ '(primes, composites, next_prime), 1825
1771           (primes, composites))                 1826
1772         (λ _ rec '(primes, composites, np), 1827
1773           rec                                         1828
1774             (if (PositiveSet.mem np composites || 1829
1775               (Z.pos np >? max))%bool%Z           1830
1776             then                                         1831
1777               (primes, composites, Pos.succ np) 1832
1778             else                                         1833
1779               (Z.pos np :: primes,            1834
1780                 List.fold_right                   1835
1781                   PositiveSet.add                1836
1782                   composites                   1837
1783                   (List.map                   1838
1784                     (λ n, Pos.mul (Pos.of_nat (S n)) np) 1839
1785                     (List.seq 0 (Z.to_nat(max/Z.pos np))), 1840
1786                     Pos.succ np)))                1841
1787               fuel                                         1842
1788               (nil, PositiveSet.empty, 2%positive))). 1843
1789
1790 Definition sieve (n : Z) := Eval cbv [sieve'] in sieve' (Z.to_nat n) n. 1844
1791
1792 We need four lemmas and an additional instance to create 1845
1793 the rewriter: 1846
1794
1795 Lemma eval_fold_right A B f x ls : 1847
1796   @List.fold_right A B f x ls 1848
1797   = ident.eagerly list_rect _ _ 1849
1798     x 1850
1799     (λ l ls fold_right_ls, f l fold_right_ls) 1851
1800     ls. 1852
1801
1802 Lemma eval_app A xs ys : 1853
1803   xs ++ ys 1854
1804   = ident.eagerly list_rect A _ 1855
1805     ys 1856
1806     (λ x xs app_xs_ys, x :: app_xs_ys) 1857
1807     xs. 1858
1808
1809 Lemma eval_map A B f ls : 1859
1810   @List.map A B f ls 1860
1811   = ident.eagerly list_rect _ _ 1861
1812     [] 1862
1813     (λ l ls map_ls, f l :: map_ls) 1863
1814     ls. 1864
1815
1816 Lemma eval_rev A xs : 1865
1817   @List.rev A xs 1866
1818   = (@list_rect _ (fun _ => _)) 1867
1819     [] 1868
1820
```

```

1871      ( $\lambda x \text{ xs} \text{ rev\_xs}, \text{ rev\_xs} ++ [x])\%list$ 
1872      xs.
1873
1874 Scheme Equality for PositiveSet.tree.
1875
1876 Definition PositiveSet_t_beq
1877   : PositiveSet.t -> PositiveSet.t -> bool
1878   := tree_beq.
1879
1880 Global Instance PositiveSet_reflect_eqb
1881   : reflect_rel (@eq PositiveSet.t) PositiveSet_t_beq
1882   := reflect_of_brel
1883     internal_tree_dec_bl internal_tree_dec_lb.
1884
1885 We then create the rewriter with
1886
1887 Make myrew := Rewriter For
1888   (eval_rect nat, eval_rect prod, eval_fold_right,
1889   eval_map, do_again eval_rev, eval_rect bool,
1890   @fst_pair, eval_rect list, eval_app)
1891   (with extra idents (Z.eqb, orb, Z.gtb,
1892     PositiveSet.elements, @fst, @snd,
1893     PositiveSet.mem, Pos.succ, PositiveSet.add,
1894     List.fold_right, List.map, List.seq, Pos.mul,
1895     S, Pos.of_nat, Z.to_nat, Z.div, Z.pos, 0,
1896     PositiveSet.empty))
1897   (with delta).
```

To get cbn and simpl to unfold our term fully, we emit

```
Global Arguments Pos.to_nat !_ / .
```

## B Additional Information on Fiat Cryptography Benchmarks

It may also be useful to see performance results with absolute times, rather than normalized execution ratios vs. the original Fiat Cryptography implementation. Furthermore, the benchmarks fit into four quite different groupings: elements of the cross product of two algorithms (unsaturated Solinas and word-by-word Montgomery) and bitwidths of target architectures (32-bit or 64-bit). Here we provide absolute-time graphs by grouping in Figure 7.

## C Experience vs. Lean and setoid\_rewrite

Although all of our toy examples work with setoid\_rewrite or rewrite\_strat (until the terms get too big), even the smallest of examples in Fiat Cryptography fell over using these tactics. When attempting to use rewrite\_strat for partial evaluation and rewriting on unsaturated Solinas with 1 limb on small primes (such as 29), we were able to get rewrite\_strat to finish after about 90 seconds. The bugs in rewrite\_strat made finding the right magic invocation quite painful, nonetheless; the invocation we settled on involved sixteen consecutive calls to rewrite\_strat with varying arguments and strategies. Trying to synthesize code for two limbs on slightly larger primes (such as 113, which needs two limbs on a 64-bit machine) took about three hours.

The widely used primes tend to have around five to ten limbs; we leave extrapolating this slowdown to the reader.

We have attached this experiment using rewrite\_strat as fiat\_crypto\_via\_rewrite\_strat.v, which is meant to be run in emacs/PG from inside the fiat-crypto directory, or in coqc by setting COQPATH to the value emitted by make printenv in fiat-crypto and then invoking the command coqc -q -R /path/to/fiat-crypto/src Crypto /path/to/fiat\_crypto\_via\_rewrite\_strat.v. To test with the two-limb prime 113, change of\_string "2^5-3" 8 in the definition of p to of\_string "2^7-15" 64.

We also tried Lean, in the hopes that rewriting in Lean, specifically optimized for performance, would be up to the challenge. Although Lean performed about 30% better than Coq on the 1-limb example, taking a bit under a minute, it did not complete on the two-limb example even after four hours (after which we stopped trying), and a five-limb example was still going after 40 hours.

We have attached our experiments with running rewrite in Lean on the Fiat Cryptography code as a supplement as well. We used Lean version 3.4.2, commit cbd2b6686ddb, Release. Run make in fiat-crypto-lean to run the one-limb example; change open ex to open ex2 to try the two-limb example, or to open ex5 to try the five-limb example.

## D Reading the Code Supplement

We have attached both the code for implementing the rewriter, as well as a copy of Fiat Cryptography adapted to use the rewriting framework. Both code supplements build with Coq 8.9 and Coq 8.10, and they require that whichever OCaml was used to build Coq be installed on the system to permit building plugins. (If Coq was installed via opam, then the correct version of OCaml will automatically be available.) Both code bases can be built by running make in the top-level directory.

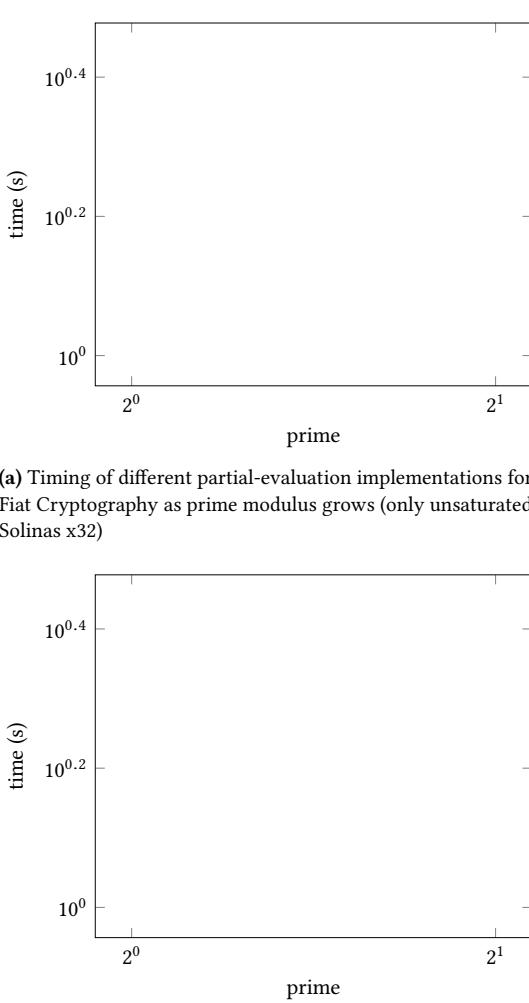
The performance data for both repositories are included at the top level as .txt and .csv files.

The performance data for the microbenchmarks can be rebuilt using make perf-SuperFast perf-Fast perf-Medium followed by make perf-csv to get the .txt and .csv files. The microbenchmarks should run in about 24 hours when run with -j5 on a 3.5 GHz machine. There also exist targets perf-Slow and perf-VerySlow, but these take significantly longer.

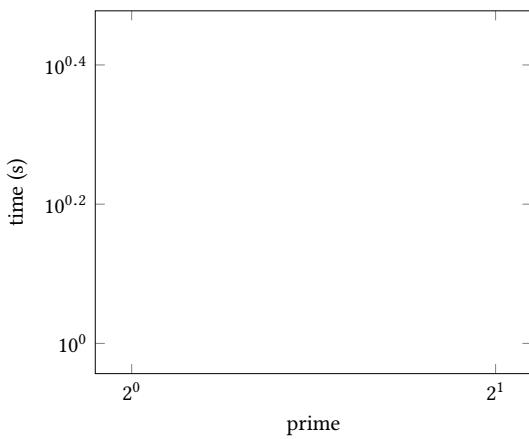
The performance data for the macrobenchmark can be rebuilt from the Fiat Cryptography copy included by running make perf -k. We ran this with PERF\_MAX\_TIME=3600 to allow each benchmark to run for up to an hour; the default is 10 minutes per benchmark. Expect the benchmarks to take over a week of time with an hour timeout and five cores. Some tests are expected to fail, making -k a necessary flag. Again, the perf-csv target will aggregate the logs and turn them into .txt and .csv files.

1926  
1927  
1928  
1929  
1930  
1931  
1932  
1933  
1934  
1935  
1936  
1937  
1938  
1939  
1940  
1941  
1942  
1943  
1944  
1945  
1946  
1947  
1948  
1949  
1950  
1951  
1952  
1953  
1954  
1955  
1956  
1957  
1958  
1959  
1960  
1961  
1962  
1963  
1964  
1965  
1966  
1967  
1968  
1969  
1970  
1971  
1972  
1973  
1974  
1975  
1976  
1977  
1978  
1979  
1980

1981  
1982  
1983  
1984  
1985  
1986  
1987  
1988  
1989  
1990  
1991  
1992  
1993  
1994  
1995  
1996  
1997  
1998  
1999  
2000  
2001  
2002  
2003  
2004  
2005  
2006  
2007  
2008  
2009  
2010  
2011  
2012  
2013  
2014  
2015  
2016  
2017  
2018  
2019  
2020  
2021  
2022  
2023  
2024  
2025  
2026  
2027  
2028  
2029  
2030  
2031  
2032  
2033  
2034  
2035

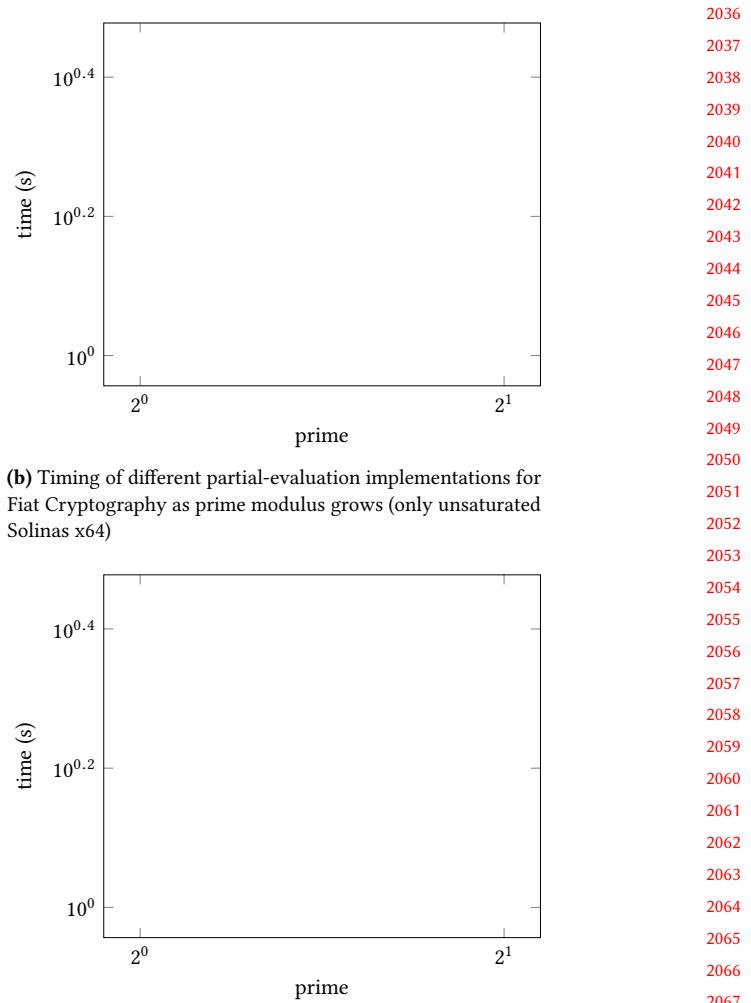


(a) Timing of different partial-evaluation implementations for Fiat Cryptography as prime modulus grows (only unsaturated Solinas x32)

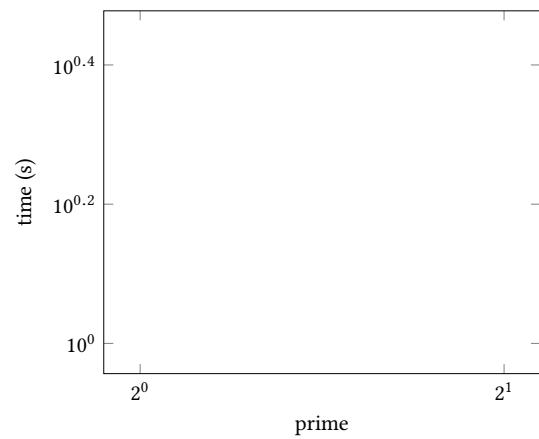


(c) Timing of different partial-evaluation implementations for Fiat Cryptography as prime modulus grows (only word-by-word Montgomery x32)

**Figure 7.** Timing of different partial-evaluation implementations for Fiat Cryptography as prime modulus grows



(b) Timing of different partial-evaluation implementations for Fiat Cryptography as prime modulus grows (only unsaturated Solinas x64)



(d) Timing of different partial-evaluation implementations for Fiat Cryptography as prime modulus grows (only word-by-word Montgomery x64)

The entry point for the rewriter is the Coq source file `rewriter/src/Rewriter/Util/plugins/RewriterBuild.v`. The rewrite rules used in Fiat Cryptography are defined in `fiat-crypto/src/Rewriter/Rules.v` and proven in `fiat-crypto/src/Rewriter/RulesProofs.v`. Note that the Fiat Cryptography copy uses COQPATH for dependency management, and `.dir-locals.el` to set COQPATH in emacs/PG; you must accept the setting when opening a file in the directory for interactive compilation to work. Thus interactive editing either requires ProofGeneral or manual setting of COQPATH. The correct value of COQPATH can be found by running `make printenv`.

We will now go through this paper and describe where to find each reference in the code base.

## D.1 Code from section 1, Introduction

### D.1.1 Code from subsection 1.1, A Motivating Example

The prefixSums example appears in the Coq source file `rewriter/src/Rewriter/Rewriter/Examples/PrefixSums.v`. Note that we use `dlet` rather than `let` in binding `acc'` so that we can preserve the `let` binder even under  $\iota$  reduction, which much of Coq's infrastructure performs eagerly. Because we attempt to isolate the dependency on the axiom of functional extensionality as much as possible, we also in practice require Proper instances for each higher-order identifier saying that each constant respects function extensionality. We hope to remove the dependency on function extensionality altogether in the future. Although we glossed over this detail in the body of this paper, we also prove

2091 **Global Instance:** `forall A B,`  
 2092   `Proper ((eq ==> eq ==> eq) ==> eq ==> eq ==> eq)`  
 2093   `(@fold_left A B).`

2094   The Make command is exposed in the file `rewriter/src/Rewriter/Util/plugins/RewriterBuild.v` and defined in  
 2095   the OCaml file `rewriter/src/Rewriter/Util/plugins/rewriter_build_plugin.mlg`. Note that one must run `make`  
 2096   to create this latter file; it is copied over from a version-  
 2097   specific file at the beginning of the build.  
 2098

2099   The `do_again`, `eval_rect`, and `ident.eagerly` constants  
 2100   are defined at the bottom of module `RewriteRuleNotations`  
 2101   in `rewriter/src/Rewriter/Language/Pre.v`.  
 2102

### 2104 D.1.2 Code from subsection 1.2, Concerns of 2105 Trusted-Code-Base Size

2106   There is no code mentioned in this section.  
 2107

### 2108 D.1.3 Code from subsection 1.3, Our Solution

2109   We claimed that our solution meets five criteria. We briefly  
 2110   justify each criterion with a sentence or a pointer to code:  
 2111

- 2112   • We claimed that we **did not grow the trusted base**  
     (excepting the axiom of functional extensionality). In  
     any example file (of which a couple can be found in  
     `rewriter/src/Rewriter/Rewriter/Examples/`), the Make command creates a rewriter package. Running `Print Assumptions` on this new constant (often named `rewriter` or `myrew`) should demonstrate a lack of axioms other than functional extensionality. `Print Assumptions` may also be run on the proof that results from using the rewriter.
- 2113   • We claimed **fast** partial evaluation with reasonable  
     memory use; we assume that the performance graphs stand on their own to support this claim. Note that memory usage can be observed by making the benchmarks while passing `TIMED=1` to `make`.
- 2114   • We claimed to allow reduction that **mixes rules of the definitional equality with equalities proven explicitly as theorems**; the “rules of the definitional equality” are, for example,  $\beta$  reduction, and we assert that it should be self-evident that our rewriter supports this.
- 2115   • We claimed common-subterm **sharing preservation**. This is implemented by supporting the use of the `dlet` notation which is defined in `rewriter/src/Rewriter/Util/LetIn.v` via the `Let_In` constant. We will come back to the infrastructure that supports this.
- 2116   • We claimed **extraction of standalone partial evaluators**. The extraction is performed in the Coq source file `perf_unsaturated_solinis.v`, in the source file `perf_word_by_word_montgomery.v`, and in the source files `saturated_solinis.v`, `unsaturated_solinis.v`, and `word_by_word_montgomery.v`, all in the directory `fiat-crypto/src/ExtractionOCaml/`. The OCaml code can be extracted and built using the target `make`

2117   `standalone-ocaml` (or `make perf-standalone` for  
 2118   the `perf_` binaries). There may be some issues with  
 2119   building these binaries on Windows as some versions  
 2120   of `ocamlopt` on Windows seem not to support out-  
 2121   putting binaries without the `.exe` extension.  
 2122

2123   The P-384 curve is mentioned. This is the curve with prime  
 2124   modulus  $2^{384} - 2^{128} - 2^6 + 2^{32} - 1$ , and the benchmarks  
 2125   for this curve can be found in the files matching the glob  
 2126   `fiat-crypto/src/Rewriter/PerfTesting/Specific/generated/p2384m2128m296p232m1_*_word_by_word_montgomery_*`.  
 2127   While the `.log` files are included in the tarball, the `.v` and  
 2128   `.sh` files are automatically generated in the course of running  
 2129   `make perf -k`.  
 2130

2131   We mention integration with abstract interpretation; the  
 2132   abstract-interpretation pass is implemented in `fiat-crypto/`  
 2133   `src/AbstractInterpretation/`.  
 2134

## 2135 D.2 Code from section 2, Trust, Reduction, and 2136 Rewriting

2137   The individual rewritings mentioned are implemented via  
 2138   the `Rewrite_*` tactics exported at the top of `rewriter/src/Rewriter/Util/plugins/RewriterBuild.v`. These tactics  
 2139   bottom out in tactics defined at the bottom of `rewriter/src/Rewriter/Rewriter/AllTactics.v`.  
 2140

### 2141 D.2.1 Code from subsection 2.1, Our Approach in 2142 Nine Steps

2143   We match the nine steps with functions from the source  
 2144   code:  
 2145

- 2146   1. The given lemma statements are scraped for which  
     named functions and types the rewriter package will  
     support. This is performed by `rewriter_scrape_data`  
     in the file `rewriter/src/Rewriter/Util/plugins/rewriter_build.ml` which invokes the tactic named  
     `make_scrape_data` in a submodule in `rewriter/src/Rewriter/Language/IdentifiersBasicGenerate.v`  
     on a goal headed by the constant we provide under the  
     name `Pre.ScrapedData.t_with_args` in `rewriter/src/Rewriter/Language/PreCommon.v`.  
 2147   2. Inductive types enumerating all available primitive  
     types and functions are emitted. This step is performed  
     by `rewriter_emit_inductives` in file `rewriter/src/Rewriter/Util/plugins/rewriter_build.ml` invoking  
     tactics, like `make_base_elim` in `rewriter/src/Rewriter/Language/IdentifiersBasicGenerate.v`,  
     on goals headed by constants from `rewriter/src/Rewriter/Language/IdentifiersBasicLibrary.v`, in-  
     cluding `base_elim_with_args` for example, to turn  
     scraped data into eliminators for the inductives. The  
     actual emitting of inductives is performed by code  
     in the file `rewriter/src/Rewriter/Util/plugins/inductive_from_elim.ml`.  
 2148

- 2201     3. Tactics generate all of the necessary definitions and  
 2202     prove all of the necessary lemmas for dealing with  
 2203     this particular set of inductive codes. This step is per-  
 2204     formed by `make_rewriter_of_scraped_and_ind` in  
 2205     the source file `rewriter/src/Rewriter/Util/plugins/`  
 2206     `rewriter_build.ml` which invokes `make_rewriter_all`  
 2207     defined in the file `rewriter/src/Rewriter/Rewriter/`  
 2208     `AllTactics.v` on a goal headed by the provided con-  
 2209     stant `VerifiedRewriter_with_ind_args` defined in  
 2210     `rewriter/src/Rewriter/Rewriter/ProofsCommon.v`.  
 2211     The definitions emitted can be found by looking at the  
 2212     tactic `Build_Rewriter` in `rewriter/src/Rewriter/`  
 2213     `Rewriter/AllTactics.v`, the tactics `build_package`  
 2214     in the source file `rewriter/src/Rewriter/Language/`  
 2215     `IdentifiersBasicGenerate.v` and also in the Coq  
 2216     source file found in `rewriter/src/Rewriter/Language/`  
 2217     `IdentifiersGenerate.v` (there is a different tactic  
 2218     named `build_package` in each of these files), and  
 2219     the tactic `prove_package_proofs_via` which can be  
 2220     found in the Coq source file `rewriter/src/Rewriter/`  
 2221     `Language/IdentifiersGenerateProofs.v`.
- 2222     4. The statements of rewrite rules are reified, and we  
 2223     prove soundness and syntactic-well-formedness lem-  
 2224     mas about each of them. This step is performed as part  
 2225     of the previous step, when the tactic `make_rewriter_all`  
 2226     transitively calls `Build_Rewriter` from `rewriter/src/`  
 2227     `Rewriter/Rewriter/AllTactics.v`. Reification is han-  
 2228     dled by the tactic `Build_RewriterT` in `rewriter/src/`  
 2229     `Rewriter/Rewriter/Reify.v`, while soundness and  
 2230     syntactic-well-formedness are handled by the tactics  
 2231     `prove_interp_good` and `prove_good` respectively, both  
 2232     in the source file `rewriter/src/Rewriter/Rewriter/`  
 2233     `ProofsCommonTactics.v`.
- 2234     5. The definitions needed to perform reification and rewrit-  
 2235     ing and the lemmas needed to prove correctness are  
 2236     assembled into a single package that can be passed  
 2237     by name to the rewriting tactic. This step is also per-  
 2238     formed by `make_rewriter_of_scraped_and_ind` in  
 2239     the source file `rewriter/src/Rewriter/Util/plugins/`  
 2240     `rewriter_build.ml`.

2241  
 2242     When we want to rewrite with a rewriter package in a  
 2243     goal, the following steps are performed, with code in the  
 2244     following places:

- 2245  
 2246  
 2247     1. We rearrange the goal into a single logical formula:  
 2248       all free-variable quantification in the proof context is  
 2249       replaced by changing the equality goal into an equal-  
 2250       ity between two functions (taking the free variables  
 2251       as inputs). Note that it is not actually an equality be-  
 2252       tween two functions but rather an `equiv` between two  
 2253       functions, where `equiv` is a custom relation we define

- 2254       indexed over type codes that is equality up to func-  
 2255       tion extensionality. This step is performed by the tac-  
 2256       tic `generalize_hyps_for_rewriting` in `rewriter/`  
 2257       `src/Rewriter/Rewriter/AllTactics.v`.
- 2258     2. We reify the side of the goal we want to simplify, using  
 2259       the inductive codes in the specified package. That side  
 2260       of the goal is then replaced with a call to a denotation  
 2261       function on the reified version. This step is performed  
 2262       by the tactic `do_reify_rhs_with` in `rewriter/src/`  
 2263       `Rewriter/Rewriter/AllTactics.v`.
- 2264     3. We use a theorem stating that rewriting preserves  
 2265       denotations of well-formed terms to replace the de-  
 2266       notation subterm with the denotation of the rewriter  
 2267       applied to the same reified term. We use Coq's built-in  
 2268       full reduction (`vm_compute`) to reduce the application  
 2269       of the rewriter to the reified term. This step is per-  
 2270       formed by the tactic `do_rewrite_with` in `rewriter/`  
 2271       `src/Rewriter/Rewriter/AllTactics.v`.
- 2272     4. Finally, we run `cbv` (a standard call-by-value reducer)  
 2273       to simplify away the invocation of the denotation  
 2274       function on the concrete syntax tree from rewriting.  
 2275       This step is performed by the tactic `do_final_cbv` in  
 2276       `rewriter/src/Rewriter/Rewriter/AllTactics.v`.

2277  
 2278     These steps are put together in the tactic `Rewrite_for_gen`  
 2279     in `rewriter/src/Rewriter/Rewriter/AllTactics.v`.

## D.2.2 Our Approach in More Than Nine Steps

2280     As the nine steps of subsection 2.1 do not exactly match  
 2281     the code, we describe here a more accurate version of what  
 2282     is going on. For ease of readability, we do not clutter this  
 2283     description with references to the code supplement, instead  
 2284     allowing the reader to match up the steps here with the more  
 2285     coarse-grained ones in subsection 2.1 or subsubsection D.2.1.

2286     In order to allow easy invocation of our rewriter, a great  
 2287     deal of code (about 6500 lines) needed to be written. Some of  
 2288     this code is about reifying rewrite rules into a form that the  
 2289     rewriter can deal with them in. Other code is about proving  
 2290     that the reified rewrite rules preserve interpretation and are  
 2291     well-formed. We wrote some plugin code to automatically  
 2292     generate the inductive type of base-type codes and identifier  
 2293     codes, as well as the two variants of the identifier-code in-  
 2294     ductive used internally in the rewriter. One interesting bit of  
 2295     code that resulted was a plugin that can emit an inductive  
 2296     declaration given the Church encoding (or eliminator) of the  
 2297     inductive type to be defined. We wrote a great deal of tactic  
 2298     code to prove basic properties about these inductive types,  
 2299     from the fact that one can unify two identifier codes and  
 2300     extract constraints on their type variables from this unifi-  
 2301     cation, to the fact that type codes have decidable equality.  
 2302     Additional plugin code was written to invoke the tactics  
 2303     that construct these definitions and prove these properties,  
 2304     so that we could generate an entire rewriter from a single

2311 command, rather than having the user separately invoke  
 2312 multiple commands in sequence.

In order to build the precomputed rewriter, the following actions are performed:

- 2315 1. The terms and types to be supported by the rewriter  
 2316 package are scraped from the given lemmas.
- 2317 2. An inductive type of codes for the types is emitted,  
 2318 and then three different versions of inductive codes for  
 2319 the identifiers are emitted (one with type arguments,  
 2320 one with type arguments supporting pattern type vari-  
 2321 ables, and one without any type arguments, to be used  
 2322 internally in pattern-matching compilation).
- 2323 3. Tactics generate all of the necessary definitions and  
 2324 prove all of the necessary lemmas for dealing with  
 2325 this particular set of inductive codes. Definitions cover  
 2326 categories like “Boolean equality on type codes” and  
 2327 “how to extract the pattern type variables from a given  
 2328 identifier code,” and lemma categories include “type  
 2329 codes have decidable equality” and “the types being  
 2330 coded for have decidable equality” and “the identifiers  
 2331 all respect function extensionality.”
- 2332 4. The rewrite rules are reified, and we prove interpretation-  
 2333 correctness and well-formedness lemmas about each  
 2334 of them.
- 2335 5. The definitions needed to perform reification and rewrit-  
 2336 ing and the lemmas needed to prove correctness are  
 2337 assembled into a single package that can be passed by  
 2338 name to the rewriting tactic.
- 2339 6. The denotation functions for type and identifier codes  
 2340 are marked for early expansion in the kernel via the  
 2341 `Strategy` command; this is necessary for conversion  
 2342 at `Qed`-time to perform reasonably on enormous goals.

When we want to rewrite with a rewriter package in a goal, the following steps are performed:

- 2344 1. We use `etransitivity` to allow rewriting separately  
 2345 on the left- and right-hand-sides of an equality. Note  
 2346 that we do not currently support rewriting in non-  
 2347 equality goals, but this is easily worked around using  
 2348 `let v := open_constr:(_) in replace <some  
 2349 term> with v` and then rewriting in the second goal.
- 2350 2. We revert all hypotheses mentioned in the goal, and  
 2351 change the form of the goal from a universally quanti-  
 2352 fied statement about equality into a statement that two  
 2353 functions are extensionally equal. Note that this step  
 2354 will fail if any hypotheses are functions not known to  
 2355 respect function extensionality via typeclass search.
- 2356 3. We reify the side of the goal that is not an existen-  
 2357 tial variable using the inductive codes in the specified  
 2358 package; the resulting goal equates the denotation of  
 2359 the newly reified term with the original evar.
- 2360 4. We use a lemma stating that rewriting preserves de-  
 2361 notations of well-formed terms to replace the goal  
 2362 with the rewriter applied to our reified term. We use

2363 `vm_compute` to prove the well-formedness side condi-  
 2364 tion reflectively. We use `vm_compute` again to reduce  
 2365 the application of the rewriter to the reified term.

- 2366 5. Finally, we run `cbv` to unfold the denotation function,  
 2367 and we instantiate the evar with the resulting rewritten  
 2368 term.

There are a couple of steps that contribute to the trusted base. We must trust that the rewriter package we generate from the rewrite rules in fact matches the rewrite rules we want to rewrite with. This involves partially trusting the scraper, the reifier, and the glue code. We must also trust the VM we use for reduction at various points in rewriting. Otherwise, everything is checked by Coq. We do, however, depend on the axiom of function extensionality in one place in the rewriter proof; after spending a couple of hours trying to remove this axiom, we temporarily gave up.

### D.3 Code from section 3, The Structure of a Rewriter

The expression language  $e$  corresponds to the inductive `expr` type defined in module `Compilers.expr` in `rewriter/src/Rewriter/Language/Language.v`.

#### D.3.1 Code from subsection 3.1, Pattern-Matching Compilation and Evaluation

The pattern -atching compilation step is done by the tactic `CompileRewrites` in `rewriter/src/Rewriter/Rewriter.Rewriter.v`, which just invokes the Gallina definition named `compile_rewrites` with ever-increasing amounts of fuel until it succeeds. (It should never fail for reasons other than insufficient fuel, unless there is a bug in the code.) The workhorse function of this code is `compile_rewrites_step`.

The decision-tree evaluation step is done by the definition `eval_rewrite_rules`, also in the file `rewriter/src/Rewriter/Rewriter.Rewriter.v`. The correctness lemmas are `eval_rewrite_rules_correct` in the file `rewriter/src/Rewriter/Rewriter/InterpProofs.v` and the theorem `wf_eval_rewrite_rules` in `rewriter/src/Rewriter/Rewriter.Wf.v`. Note that the second of these lemmas, not mentioned in the paper, is effectively saying that for two related syntax trees, `eval_rewrite_rules` picks the same rewrite rule for both. (We actually prove a slightly weaker lemma, which is a bit harder to state in English.)

The third step of rewriting with a given rule is performed by the definition `rewrite_with_rule` in `rewriter/src/Rewriter/Rewriter.Rewriter.Rewriter.v`. The correctness proof is `interp_rewrite_with_rule` in `rewriter/src/Rewriter/Rewriter/InterpProofs.v`. Note that the well-formedness-preservation proof for this definition is inlined into the proof `wf_eval_rewrite_rules` mentioned above.

The inductive description of decision trees is `decision_tree` in `rewriter/src/Rewriter/Rewriter.Rewriter.v`.

2369  
 2370  
 2371  
 2372  
 2373  
 2374  
 2375  
 2376  
 2377  
 2378  
 2379  
 2380  
 2381  
 2382  
 2383  
 2384  
 2385  
 2386  
 2387  
 2388  
 2389  
 2390  
 2391  
 2392  
 2393  
 2394  
 2395  
 2396  
 2397  
 2398  
 2399  
 2400  
 2401  
 2402  
 2403  
 2404  
 2405  
 2406  
 2407  
 2408  
 2409  
 2410  
 2411  
 2412  
 2413  
 2414  
 2415  
 2416  
 2417  
 2418  
 2419  
 2420

2421     The pattern language is defined as the inductive pattern  
 2422     in `rewriter/src/Rewriter/Rewriter/Rewriter.v`. Note  
 2423     that we have a Raw version and a typed version; the pattern-  
 2424     matching compilation and decision-tree evaluation of Aehlig  
 2425     et al. [1] is an algorithm on untyped patterns and untyped  
 2426     terms. We found that trying to maintain typing constraints  
 2427     led to headaches with dependent types. Therefore when  
 2428     doing the actual decision-tree evaluation, we wrap all of our  
 2429     expressions in the dynamically typed `rawexpr` type and all  
 2430     of our patterns in the dynamically typed `Raw.pattern` type.  
 2431     We also emit separate inductives of identifier codes for each  
 2432     of the `expr`, `pattern`, and `Raw.pattern` type families.

2433     We partially evaluate the partial evaluator defined by  
 2434     `eval_rewrite_rules` in the tactic `make_rewrite_head` in  
 2435     `rewriter/src/Rewriter/Rewriter/Reify.v`.

### 2437 D.3.2 Code from subsection 3.2, Adding 2438     Higher-Order Features

2439     The type `NbEt` mentioned in this paper is not actually used in  
 2440     the code; the version we have is described in subsection 4.2 as  
 2441     the definition `value'` in `rewriter/src/Rewriter/Rewriter/`  
 2442     `Rewriter.v`.

2443     The functions `reify` and `reflect` are defined in `rewriter/`  
 2444     `src/Rewriter/Rewriter/Rewriter.v` and share names with  
 2445     the functions in the paper. The function `reduce` is named  
 2446     `rewrite_bottomup` in the code, and the closest match to  
 2447     `NbE` is `rewrite`.

## 2449 D.4 Code from section 4, Scaling Challenges

### 2451 D.4.1 Code from subsection 4.1, Variable 2452     Environments Will Be Large

2453     The inductives type, `base_type` (actually the inductive type  
 2454     `base.type.type` in the supplemental code), and `expr`, as  
 2455     well as the definition `Expr`, are all defined in `rewriter/src/`  
 2456     `Rewriter/Language/Language.v`. The definition `denoteT`  
 2457     is the fixpoint type.`.interp` (the fixpoint `interp` in the  
 2458     module type) in `rewriter/src/Rewriter/Language/Language.v`.  
 2459     The definition `denoteE` is `expr.interp`, and `DenoteE` is the  
 2460     fixpoint `expr.interp`.

2461     As mentioned above, `nbeT` does not actually exist as stated  
 2462     but is close to `value'` in `rewriter/src/Rewriter/Rewriter/`  
 2463     `Rewriter.v`. The functions `reify` and `reflect` are defined  
 2464     in `rewriter/src/Rewriter/Rewriter/Rewriter.v` and share  
 2465     names with the functions in the paper. The actual code is  
 2466     somewhat more complicated than the version presented  
 2467     in the paper, due to needing to deal with converting well-  
 2468     typed-by-construction expressions to dynamically typed ex-  
 2469     pressions for use in decision-tree evaluation and also due  
 2470     to the need to support early partial evaluation against a  
 2471     concrete decision tree. Thus the version of `reflect` that  
 2472     actually invokes rewriting at base types is a separate defi-  
 2473     nition `assemble_identifier_rewriters`, while `reify` in-  
 2474     volves a version of `reflect` (named `reflect`) that does not

2475     call rewriting. The function named `reduce` is what we call  
 2476     `rewrite_bottomup` in the code; the name `Rewrite` is shared  
 2477     between this paper and the code. Note that we eventually in-  
 2478     stantiate the argument `rewrite_head` of `rewrite_bottomup`  
 2479     with a partially evaluated version of the definition named  
 2480     `assemble_identifier_rewriters`. Note also that we use  
 2481     `fuel` to support `do_again`, and this is used in the definition  
 2482     `repeat_rewrite` that calls `rewrite_bottomup`.

2483     The correctness theorems are `InterpRewrite` in `rewriter/`  
 2484     `src/Rewriter/Rewriter/InterpProofs.v` and `Wf_Rewrite`  
 2485     in `rewriter/src/Rewriter/Rewriter/Wf.v`.

2486     Packages containing rewriters and their correctness the-  
 2487     orems are in the record `VerifiedRewriter` in `rewriter/`  
 2488     `src/Rewriter/Rewriter/ProofsCommon.v`; a package of  
 2489     this type is then passed to the tactic `Rewrite_for_gen` from  
 2490     `rewriter/src/Rewriter/Rewriter/AllTactics.v` to per-  
 2491     form the actual rewriting. The correspondence of the code  
 2492     to the various steps in rewriting is described in the second  
 2493     list of subsubsection D.2.1.

### 2494 D.4.2 Code from subsection 4.2, Subterm Sharing is 2495     Crucial

2496     To run the P-256 example in the copy of Fiat Cryptography  
 2497     attached as a code supplement, after building the library, run  
 2498     the code

```
2499 Require Import Crypto.Rewriter.PerfTesting.Core.  
2500 Require Import Crypto.Util.Option.
```

```
2501 Import WordByWordMontgomery.  
2502 Import Core.RuntimeDefinitions.
```

```
2503 Definition p : params  
2504   := Eval compute in invert_Some  
2505   (of_string "2^256-2^224+2^192+2^96-1" 64).
```

```
2506 Goal True.  
2507 (* Successful run: *)  
2508 Time let v := (eval cbv  
2509   -[Let_In  
2510     runtime_nth_default  
2511     runtime_add  
2512     runtime_sub  
2513     runtime_mul  
2514     runtime_opp  
2515     runtime_div  
2516     runtime_modulo  
2517     RT_Z.add_get_carry_full  
2518     RT_Z.add_with_get_carry_full  
2519     RT_Z.mul_split]  
2520   in (GallinaDefOf p)) in  
2521   idtac.
```

```
2522 (* Unsuccessful OOM run: *)  
2523 Time let v := (eval cbv  
2524   -[(*)Let_In*)  
2525     runtime_nth_default  
2526     runtime_add  
2527     runtime_sub  
2528     runtime_mul  
2529     runtime_opp  
2530     runtime_div  
2531     runtime_modulo  
2532     RT_Z.add_get_carry_full  
2533     RT_Z.add_with_get_carry_full  
2534     RT_Z.mul_split]  
2535   in (GallinaDefOf p)) in  
2536   idtac.
```

```

2531     runtime_sub
2532     runtime_mul
2533     runtime_opp
2534     runtime_div
2535     runtime_modulo
2536     RT_Z.add_get_carry_full
2537     RT_Z.add_with_get_carry_full
2538     RT_Z.mul_split]
2539   in (GallinaDefOf p)) in
2540   idtac.
2541 Abort.

2542 The UnderLets monad is defined in the file rewriter/
2543 src/Rewriter/Language/UnderLets.v.
2544 The definitions nbeT', nbeT, and nbeT_with_lets are in
2545 rewriter/src/Rewriter/Rewriter/Rewriter.v and are
2546 named value', value, and value_with_lets, respectively.

2547
2548 D.4.3 Code from subsection 4.3, Rules Need Side
2549 Conditions
2550 The “variant of pattern variable that only matches constants”
2551 is actually special support for the reification of ident.literal
2552 (defined in the module RewriteRuleNotations in rewriter/
2553 src/Rewriter/Language/Pre.v) threaded throughout the
2554 rewriter. The apostrophe notation ' is also introduced in the
2555 module RewriteRuleNotations in rewriter/src/Rewriter/
2556 Language/Pre.v. The support for side conditions is handled
2557 by permitting rewrite-rule-replacement expressions to re-
2558 turn option expr instead of expr, allowing the function
2559 expr_to_pattern_and_replacement in the file rewriter/
2560 src/Rewriter/Rewriter/Reify.v to fold the side condi-
2561 tions into a choice of whether to return Some or None.
2562
2563 D.4.4 Code from subsection 4.4, Side Conditions
2564 Need Abstract Interpretation
2565 The abstract-interpretation pass is defined in fiat-crypto/
2566 src/AbstractInterpretation/, and the rewrite rules han-
2567 dling abstract-interpretation results are the Gallina defi-
2568 nitions arith_with_casts_rewrite_rulesT, in addition
2569 to strip_literal_casts_rewrite_rulesT, in addition to
2570 fancy_with_casts_rewrite_rulesT, and finally in addi-
2571 tion to mul_split_rewrite_rulesT, all defined in fiat-crypto/
2572 src/Rewriter/Rules.v.
2573 The clip function is the definition ident.cast in fiat-crypto/
2574 src/Language/PreExtra.v.
2575
2576 D.5 Code from section 5, Evaluation
2577 D.5.1 Code from subsection 5.1, Microbenchmarks
2578 This code is found in the files in rewriter/src/Rewriter/
2579 Rewriter/Examples/. We ran the microbenchmarks using
2580 the code in rewriter/src/Rewriter/Rewriter/Examples/
2581 PerfTesting/Harness.v together with some Makefile clev-
2582 erness. The file names correspond to the section titles in
2583 Appendix A.
2584

```

<b>D.5.2 Code from subsection 5.2, Macrobenchmark:</b>	2586
<b>Fiat Cryptography</b>	2587
The rewrite rules are defined in fiat-crypto/src/Rewriter/	
Rules.v and proven in the file fiat-crypto/src/Rewriter/	
RulesProofs.v. They are turned into rewriters in the vari-	
ous files in fiat-crypto/src/Rewriter/Passes/. The shared	
inductives and definitions are defined in the Coq source files	
fiat-crypto/src/Language/IdentifiersBasicGENERATED.v,	
fiat-crypto/src/Language/IdentifiersGENERATED.v, and	
fiat-crypto/src/Language/IdentifiersGENERATEDProofs.v.	
Note that we invoke the subtactics of the Make command	
manually to increase parallelism in the build and to allow a	
shared language across multiple rewriter packages.	